

**Advanced Modular Cryptocurrency Trading
&
Analytics Platform:
A Comprehensive Mathematical and
Implementation Framework**

Sineshaw (Shaw) Mesfin Tesfaye

Department of Computer Science

Georgia State University

`stesfaye4@student.gsu.edu`

2025

Cryptocurrency Trading Research — Integrated Submission

Methods Transparency Checklist

Category	Details
Data & Instruments	Coinbase Advanced Trade API; BTC-USD, ETH-USD, SOL-USD, and 20+ altcoin pairs; 1-min to daily granularity; Jan 2023–Dec 2024
Targets & Horizons	1–4 hour price direction classification; next-candle price regression; multi-horizon forecasting (5-min, 1-hr, 4-hr)
Cross-Validation	Walk-forward validation with expanding window; 70/15/15 train/validation/test split; purged k -fold to prevent look-ahead bias
Costs & Slippage	0.6% round-trip transaction cost (Coinbase Advanced); 10 bps slippage model; market impact via square-root model
Risk Controls	Maximum 25% portfolio risk per position; Kelly criterion position sizing with half-Kelly adjustment; GARCH-based stop-loss calibration

Abstract

This paper presents a comprehensive, modular cryptocurrency trading and analytics platform that integrates advanced machine learning, quantitative finance, and real-time market data processing. The system employs a stacking ensemble of Random Forest, XGBoost, Support Vector Regression, Ridge, and Lasso regressors with a gradient boosting meta-learner, achieving directional accuracy exceeding 62% on out-of-sample cryptocurrency data. We develop 35+ WorldQuant-style alpha factors spanning price-based, momentum, cross-sectional, and statistical categories. Risk management is formalized through GARCH(1,1) volatility modeling, Value at Risk with Expected Shortfall, Kelly criterion position sizing, and factor model risk decomposition. Signal denoising via Kalman filtering and wavelet transforms improves feature quality. Walk-forward validation and Kupiec backtesting confirm model robustness across bull, bear, and sideways market regimes. The platform processes live WebSocket feeds from Coinbase Advanced Trade API and executes trades through an event-driven architecture with sub-second latency. All mathematical foundations, implementation details, and empirical results are documented for full reproducibility.

Keywords: cryptocurrency trading, machine learning, ensemble methods, alpha factors, risk management, GARCH, Kelly criterion, stacking ensemble, quantitative finance, Coinbase API

Reproducibility Statement

All source code, configuration files, and trained model artifacts are available in the project repository. Data is sourced exclusively from public Coinbase API endpoints. Random seeds are fixed for all stochastic components. Database schemas and SQL queries are provided in the appendices.

Ethics Statement

This research involves automated trading of publicly listed cryptocurrency assets. No insider information or non-public data sources are used. The system includes circuit breakers, maximum position limits, and risk controls to prevent excessive market impact. All trading is conducted on the author's personal accounts.

Limitations

Results are based on historical and live simulation data from a single exchange (Coinbase). Cryptocurrency markets exhibit regime changes, structural breaks, and liquidity variations that may cause future performance to differ materially from reported results.

Transaction costs and slippage are modeled but may underestimate real-world execution friction during high-volatility periods.

Contents

1	Introduction	9
1.1	Motivation	9
1.2	Key Contributions	9
1.3	System Overview	9
2	Background and Literature Review	10
2.1	Financial Machine Learning	10
2.2	Ensemble Methods in Finance	10
2.3	Cryptocurrency Market Research	10
2.4	Technical Analysis and Market Microstructure	11
2.5	Deep Learning for Time Series	11
2.6	Risk Management	11
3	Mathematical Framework	11
3.1	Technical Indicators	11
3.1.1	Exponential Moving Average (EMA)	11
3.1.2	Relative Strength Index (RSI)	12
3.1.3	Moving Average Convergence Divergence (MACD)	12
3.1.4	Bollinger Bands	12
3.1.5	Stochastic Oscillator	12
3.1.6	Average True Range (ATR)	13
3.1.7	On-Balance Volume (OBV)	13
3.2	Advanced Statistical Features	13
3.2.1	Rolling Volatility	13
3.2.2	Maximum Drawdown	13
3.2.3	Sharpe Ratio	13
3.2.4	Sortino Ratio	13
3.2.5	Calmar Ratio	13
3.3	WorldQuant-Style Alpha Factors	13
3.3.1	Price-Based Factors (Factors 1–10)	14
3.3.2	Momentum Factors (Factors 11–20)	15
3.3.3	Cross-Sectional Factors (Factors 21–28)	15
3.3.4	Statistical Factors (Factors 29–35+)	16
3.4	Machine Learning Models	16
3.4.1	Random Forest	16
3.4.2	Gradient Boosting and XGBoost	17
3.4.3	Ridge and Lasso Regression	18
3.4.4	Support Vector Regression (SVR)	18

3.4.5	LSTM Networks	19
3.4.6	Stacking Ensemble	19
3.5	Portfolio Optimization and Weight Allocation	20
3.5.1	Minimum Variance Portfolio	20
3.5.2	Mean-Variance Optimization	20
3.5.3	Bayesian Model Averaging	20
3.6	Risk Management	20
3.6.1	GARCH(1,1) Volatility Model	20
3.6.2	Value at Risk (VaR)	21
3.6.3	Expected Shortfall (Conditional VaR)	21
3.6.4	Kelly Criterion Position Sizing	22
3.6.5	Factor Model Risk Decomposition	23
3.7	Signal Denoising	23
3.7.1	Kalman Filter	23
3.7.2	Wavelet Transform Denoising	24
3.8	Volatility Clustering and Regime Detection	24
3.8.1	Volatility Clustering Detection	24
3.8.2	Mean Reversion Testing	24
4	System Architecture	25
4.1	Modular Design	25
4.2	Key Components	25
4.3	Data Flow	25
5	Machine Learning Methodology	26
5.1	Feature Engineering Pipeline	26
5.2	Model Selection and Training	26
5.3	Production Feature Set	27
5.4	Memory-Efficient Processing	27
6	Risk Management Framework	27
6.1	Multi-Layer Risk Controls	27
6.2	Dynamic Position Sizing	28
6.3	Execution and Order Management	28
6.4	Alpha Factor Integration	28
7	User Experience and Analytics	28
7.1	Trading Dashboard	28
7.2	Analytics and Reporting	29
7.3	Backtesting Framework	29

8 Performance Analysis	29
8.1 Classification Performance (Entry Signals)	29
8.2 Cross-Asset Generalisation (60 Symbols)	31
8.3 Cross-Validation Performance	33
8.4 Feature Category Performance	33
8.5 System Design Rationale: TP/SL Exit Management	34
8.6 Statistical Testing	34
8.7 Production System Performance	35
8.8 Quantitative Backtesting Results	35
8.8.1 Strategy Comparison	36
8.8.2 Per-Asset Breakdown	37
8.8.3 Trade-Level Analysis	38
8.8.4 Summary of Backtesting Findings	39
9 Validation and Robustness	40
9.1 Walk-Forward Validation	40
9.2 VaR Backtesting: Kupiec Test	40
9.3 Factor Stability Analysis	40
9.4 Christoffersen Conditional Coverage Test	41
10 Discussion	41
10.1 Key Findings	41
10.2 Comparison with Literature	41
10.3 Computational Considerations	41
10.4 Limitations and Challenges	42
10.5 Practical Applications	42
11 Conclusion	42
11.1 Summary	42
11.2 Implications	43
11.3 Future Research	43
A System Requirements	44
B Configuration Parameters	45
C Complete Alpha Factor Catalog	45
C.1 Price-Based Alpha Factors	45
C.2 Momentum Alpha Factors	46
C.3 Cross-Sectional Alpha Factors	46
C.4 Statistical Alpha Factors	47

C.5	High-Frequency Factors	47
D	Kalman Filter Implementation	47
E	Feature Selection Algorithm	48
F	Database Schema	48
G	Stacking Ensemble Approach for Trading Signals	49
G.1	Abstract	50
G.2	Base Model Architecture	50
G.3	Meta-Learning Framework	50
G.4	Signal Generation	50
H	Mathematical Foundations of Base Models	51
H.1	Ridge Regression: Derivation	51
H.2	Lasso Regression: Coordinate Descent	51
H.3	SVR: Dual Formulation	52
H.4	Gradient Boosting: Functional Gradient Descent	52
H.5	XGBoost: Regularized Objective	52

1 Introduction

1.1 Motivation

Cryptocurrency markets present unique challenges and opportunities for algorithmic trading systems. Unlike traditional equity markets, crypto assets trade 24/7 across fragmented exchanges with varying liquidity, exhibit extreme volatility, and are influenced by sentiment-driven dynamics that differ fundamentally from established financial instruments (??). These characteristics demand trading systems that combine rigorous quantitative methods with robust engineering.

1.2 Key Contributions

This work makes the following contributions:

- (1) A **modular system architecture** with cleanly separated components for data ingestion, feature engineering, model training, risk management, order execution, and analytics.
- (2) **35+ WorldQuant-style alpha factors** organized into price-based, momentum, cross-sectional, and statistical categories, with formal mathematical definitions.
- (3) A **stacking ensemble** combining Random Forest, XGBoost, SVR, Ridge, and Lasso base learners with a gradient boosting meta-learner.
- (4) **Comprehensive risk management** integrating GARCH(1,1) volatility forecasting, parametric and historical VaR, Expected Shortfall, Kelly criterion position sizing, and factor model risk decomposition.
- (5) **Signal denoising** via Kalman filtering and wavelet transforms for improved feature quality.
- (6) **Rigorous validation** through walk-forward testing, Kupiec VaR backtesting, bootstrap confidence intervals, and regime-conditional performance analysis.

1.3 System Overview

The platform is implemented in Python and consists of the following primary modules:

- `maybe.py` — Core trading logic, indicator calculation, model training, and live execution.
- `flask_trading_dashboard.py` — Web-based dashboard with real-time analytics, GARCH-based TP/SL, and backtesting.

- `stefan_jansen_improvements.py` — Enhanced feature engineering with alpha factors, momentum features, and denoising.
- `quantitative_finance_ml_enhancement.py` — GARCH, regime detection, Kelly criterion, and Monte Carlo simulation.
- `portfolio_var.py` — Portfolio VaR, Component VaR, and Expected Shortfall calculations.
- `factor_risk.py` — Factor model regression and risk decomposition.
- `enhanced_features.py` — Technical indicator library (RSI, Bollinger Bands, volume ratios).
- `stacking_ml_engine.py` — Stacking ensemble decision engine.

2 Background and Literature Review

2.1 Financial Machine Learning

The application of machine learning to financial markets has grown substantially since the seminal work of ? on the Adaptive Markets Hypothesis. ? provides a comprehensive treatment of ML techniques for algorithmic trading, including factor research, ensemble methods, and alternative data integration. ? introduces innovations such as the triple barrier method, fractional differentiation, and purged cross-validation that address unique challenges in financial ML.

2.2 Ensemble Methods in Finance

Ensemble learning combines multiple base models to improve prediction accuracy and robustness. ? introduced Random Forests as bagged decision tree ensembles. ? developed XGBoost, a scalable gradient boosting framework that has become dominant in structured data prediction. ? formalized stacked generalization, where a meta-learner combines base model predictions. In the cryptocurrency context, ? demonstrates that stacking ensembles outperform individual models for trading signal generation.

2.3 Cryptocurrency Market Research

Cryptocurrency markets are characterized by high volatility, 24/7 trading, fragmented liquidity, and susceptibility to sentiment-driven price movements (??). ? document that standard risk factors fail to explain cryptocurrency returns, motivating the development of crypto-specific alpha factors. ? find that momentum and reversal effects exist in crypto markets but differ in magnitude and persistence from equity markets.

2.4 Technical Analysis and Market Microstructure

Technical analysis remains widely used in cryptocurrency trading despite mixed evidence of its efficacy in traditional markets (?). The key indicators implemented in our system include exponential moving averages, relative strength index, MACD, Bollinger Bands, and stochastic oscillators. Market microstructure theory (?) informs our modeling of price impact, bid-ask spreads, and liquidity.

2.5 Deep Learning for Time Series

Long Short-Term Memory (LSTM) networks (?) have shown promise for financial time series prediction due to their ability to capture long-range dependencies. ? demonstrate that LSTMs outperform traditional models for stock return prediction over certain horizons.

2.6 Risk Management

Modern portfolio risk management builds on the work of ? and has been extended through Value at Risk (?), Expected Shortfall (?), and the Kelly criterion (??). GARCH models (??) remain the standard for volatility forecasting in financial applications.

3 Mathematical Framework

3.1 Technical Indicators

3.1.1 Exponential Moving Average (EMA)

The exponential moving average assigns exponentially decreasing weights to past observations. For a price series $\{P_t\}$ with span parameter n :

$$\text{EMA}_t = \alpha \cdot P_t + (1 - \alpha) \cdot \text{EMA}_{t-1}, \quad \alpha = \frac{2}{n+1}, \quad (1)$$

where α is the smoothing factor and $\text{EMA}_0 = P_0$. The system computes EMA_{12} and EMA_{26} as implemented in `maybe.py`:

```
1 df['EMA12'] = df['close'].ewm(span=12, adjust=False).mean()
2 df['EMA26'] = df['close'].ewm(span=26, adjust=False).mean()
```

Listing 1: EMA calculation in `maybe.py`

3.1.2 Relative Strength Index (RSI)

The RSI measures the magnitude of recent price changes to evaluate overbought or oversold conditions:

$$\text{RSI}_t = 100 - \frac{100}{1 + \text{RS}_t}, \quad \text{RS}_t = \frac{\overline{\text{Gain}}_t}{\overline{\text{Loss}}_t}, \quad (2)$$

where

$$\overline{\text{Gain}}_t = \frac{1}{n} \sum_{i=0}^{n-1} \max(\Delta P_{t-i}, 0), \quad (3)$$

$$\overline{\text{Loss}}_t = \frac{1}{n} \sum_{i=0}^{n-1} \max(-\Delta P_{t-i}, 0), \quad (4)$$

with $\Delta P_t = P_t - P_{t-1}$ and $n = 14$ by default.

3.1.3 Moving Average Convergence Divergence (MACD)

$$\text{MACD}_t = \text{EMA}_{12,t} - \text{EMA}_{26,t}, \quad (5)$$

$$\text{Signal}_t = \text{EMA}_9(\text{MACD}_t), \quad (6)$$

$$\text{Histogram}_t = \text{MACD}_t - \text{Signal}_t. \quad (7)$$

3.1.4 Bollinger Bands

Bollinger Bands construct an envelope around a moving average using standard deviation:

$$\text{Middle}_t = \text{SMA}_{20,t} = \frac{1}{20} \sum_{i=0}^{19} P_{t-i}, \quad (8)$$

$$\text{Upper}_t = \text{SMA}_{20,t} + k \cdot \sigma_{20,t}, \quad (9)$$

$$\text{Lower}_t = \text{SMA}_{20,t} - k \cdot \sigma_{20,t}, \quad (10)$$

where $\sigma_{20,t}$ is the rolling 20-period standard deviation and $k = 2$ by default. The Bollinger Band position indicator normalizes the current price within the band:

$$\text{BB_Position}_t = \frac{P_t - \text{Lower}_t}{\text{Upper}_t - \text{Lower}_t}. \quad (11)$$

3.1.5 Stochastic Oscillator

$$\%K_t = 100 \cdot \frac{P_t - L_{14,t}}{H_{14,t} - L_{14,t}}, \quad (12)$$

$$\%D_t = \text{SMA}_3(\%K_t), \quad (13)$$

where $H_{14,t} = \max_{i \in [t-13,t]} H_i$ and $L_{14,t} = \min_{i \in [t-13,t]} L_i$.

3.1.6 Average True Range (ATR)

$$\text{ATR}_t = \frac{1}{14} \sum_{i=0}^{13} \text{TR}_{t-i}, \quad \text{TR}_t = \max(H_t - L_t, |H_t - C_{t-1}|, |L_t - C_{t-1}|). \quad (14)$$

3.1.7 On-Balance Volume (OBV)

$$\text{OBV}_t = \text{OBV}_{t-1} + \text{sign}(\Delta P_t) \cdot V_t. \quad (15)$$

3.2 Advanced Statistical Features

3.2.1 Rolling Volatility

Annualized rolling volatility over window w :

$$\sigma_{w,t} = \sqrt{\frac{252}{w} \sum_{i=0}^{w-1} (r_{t-i} - \bar{r}_{w,t})^2}, \quad r_t = \ln \frac{P_t}{P_{t-1}}. \quad (16)$$

3.2.2 Maximum Drawdown

$$\text{MDD}_t = \max_{s \in [0,t]} \frac{M_s - P_s}{M_s}, \quad M_s = \max_{u \in [0,s]} P_u. \quad (17)$$

3.2.3 Sharpe Ratio

$$\text{SR} = \frac{\bar{r} - r_f}{\sigma_r} \cdot \sqrt{252}, \quad (18)$$

where r_f is the risk-free rate and σ_r is the standard deviation of returns.

3.2.4 Sortino Ratio

$$\text{Sortino} = \frac{\bar{r} - r_f}{\sigma_d} \cdot \sqrt{252}, \quad \sigma_d = \sqrt{\frac{1}{N} \sum_{i=1}^N \min(r_i - r_f, 0)^2}. \quad (19)$$

3.2.5 Calmar Ratio

$$\text{Calmar} = \frac{\text{CAGR}}{\text{MDD}}. \quad (20)$$

3.3 WorldQuant-Style Alpha Factors

We implement 35+ alpha factors organized into four categories. Each factor is standardized via cross-sectional z -scoring or rank normalization.

3.3.1 Price-Based Factors (Factors 1–10)

$$\alpha_1 = \text{rank} \left(\arg \max_{i \in [1, \text{ts}]} \text{SignedPower}(\text{returns}, 2) \right) - 0.5, \quad (21)$$

$$\alpha_2 = -1 \times \text{correlation}(\text{rank}(\delta(\log(V_t), 2)), \text{rank} \left(\frac{C_t - O_t}{H_t - L_t} \right), 6), \quad (22)$$

$$\alpha_3 = -1 \times \text{correlation}(\text{rank}(O_t), \text{rank}(V_t), 10), \quad (23)$$

$$\alpha_4 = -1 \times \text{Ts_Rank}(\text{rank}(C_t), 9), \quad (24)$$

$$\alpha_5 = \text{rank}(O_t - \text{Ts_Min}(O_t, 12)) \times [\text{rank}(\text{correlation}(V_t, \text{SMA}_5(V_t), 26))]^5, \quad (25)$$

$$\alpha_6 = -1 \times \text{correlation}(O_t, V_t, 10), \quad (26)$$

$$\alpha_7 = \begin{cases} -1 \times \text{Ts_Rank}(|r_t|, 5) & \text{if } \text{ADV}_{20} < V_t, \\ -1 & \text{otherwise,} \end{cases} \quad (27)$$

$$\alpha_8 = -1 \times \text{rank} \left(\delta \left(\frac{(C_t - L_t) - (H_t - C_t)}{H_t - L_t} \times V_t, 1 \right) \right), \quad (28)$$

$$\alpha_9 = \begin{cases} \delta(C_t, 1) \times (C_t - \text{Ts_Min}(C_t, 5)) & \text{if } \min(\delta(C_t, 1), 0) > 0, \\ \delta(C_t, 1) & \text{otherwise,} \end{cases} \quad (29)$$

$$\alpha_{10} = \text{rank} \left(\begin{cases} \delta(C_t, 1) & \text{if } \min(\delta(C_t, 1), 0) > 0, \\ \delta(C_t, 1) & \text{otherwise} \end{cases} \right), \quad (30)$$

where $\delta(X_t, d) = X_t - X_{t-d}$, $\text{Ts_Rank}(X, d)$ is the time-series rank of X over the past d periods, $\text{Ts_Min}(X, d) = \min_{i \in [t-d+1, t]} X_i$, and ADV_{20} is the 20-period average daily volume.

3.3.2 Momentum Factors (Factors 11–20)

$$\alpha_{11} = \frac{C_t}{C_{t-5}} - 1, \quad (31)$$

$$\alpha_{12} = \frac{C_t}{C_{t-20}} - 1, \quad (32)$$

$$\alpha_{13} = \frac{C_t}{C_{t-60}} - 1, \quad (33)$$

$$\alpha_{14} = \frac{C_t}{C_{t-120}} - 1, \quad (34)$$

$$\alpha_{15} = \frac{C_t}{C_{t-252}} - 1, \quad (35)$$

$$\alpha_{16} = \frac{\partial P}{\partial t} \approx \Delta P_t = P_t - P_{t-1} \quad (\text{price velocity}), \quad (36)$$

$$\alpha_{17} = \frac{\partial^2 P}{\partial t^2} \approx \Delta^2 P_t = \Delta P_t - \Delta P_{t-1} \quad (\text{price acceleration}), \quad (37)$$

$$\alpha_{18} = \frac{\partial^3 P}{\partial t^3} \approx \Delta^3 P_t \quad (\text{price jerk}), \quad (38)$$

$$\alpha_{19} = \text{VPT}_t = \sum_{i=1}^t V_i \cdot r_i \quad (\text{Volume-Price Trend}), \quad (39)$$

$$\alpha_{20} = \frac{\text{VPT}_t - \text{VPT}_{t-10}}{\text{VPT}_{t-10}} \quad (\text{VPT Momentum}). \quad (40)$$

3.3.3 Cross-Sectional Factors (Factors 21–28)

$$\alpha_{21} = \text{AD}_t = \sum_{i=1}^t \frac{(C_i - L_i) - (H_i - C_i)}{H_i - L_i} \cdot V_i \quad (\text{A/D Line}), \quad (41)$$

$$\alpha_{22} = \frac{\text{AD}_t - \text{AD}_{t-10}}{\text{AD}_{t-10}} \quad (\text{A/D Momentum}), \quad (42)$$

$$\alpha_{23} = \text{OBV}_t = \sum_{i=1}^t \text{sign}(r_i) \cdot V_i, \quad (43)$$

$$\alpha_{24} = \frac{\text{OBV}_t - \text{OBV}_{t-10}}{\text{OBV}_{t-10}} \quad (\text{OBV Momentum}), \quad (44)$$

$$\alpha_{25} = \mathbb{1} \left[C_t > \max_{i \in [t-d, t-1]} H_i \right] \quad (\text{Breakout High, } d \in \{10, 20, 50\}), \quad (45)$$

$$\alpha_{26} = \mathbb{1} \left[C_t < \min_{i \in [t-d, t-1]} L_i \right] \quad (\text{Breakout Low, } d \in \{10, 20, 50\}), \quad (46)$$

$$\alpha_{27} = \frac{|\text{SMA}_{10,t} - \text{SMA}_{20,t}|}{C_t} \quad (\text{Trend Strength Short}), \quad (47)$$

$$\alpha_{28} = \frac{|\text{SMA}_{20,t} - \text{SMA}_{50,t}|}{C_t} \quad (\text{Trend Strength Long}). \quad (48)$$

3.3.4 Statistical Factors (Factors 29–35+)

$$\alpha_{29} = \frac{C_t}{\text{SMA}_{d,t}} - 1 \quad (\text{Relative Strength, } d \in \{10, 20, 50\}), \quad (49)$$

$$\alpha_{30} = \frac{\bar{r}_{d,t}}{\sigma_{d,t}} \quad (\text{Sharpe-like Momentum, } d \in \{10, 20\}), \quad (50)$$

$$\alpha_{31} = \frac{V_t}{\text{SMA}_{20}(V_t)} \quad (\text{Volume Ratio}), \quad (51)$$

$$\alpha_{32} = \text{rank}(\sigma_{10,t}) \quad (\text{Volatility Rank}), \quad (52)$$

$$\alpha_{33} = \text{rank}(\text{skew}_{20,t}) \quad (\text{Skewness Rank}), \quad (53)$$

$$\alpha_{34} = \text{rank}(\text{kurt}_{20,t}) \quad (\text{Kurtosis Rank}), \quad (54)$$

$$\alpha_{35} = H(r_{t-w:t}) \quad (\text{Hurst Exponent}), \quad (55)$$

where the Hurst exponent is estimated via the rescaled range statistic:

$$H = \frac{\log(R/S)}{\log(n)}, \quad \frac{R}{S} = \frac{\max_{1 \leq k \leq n} W_k - \min_{1 \leq k \leq n} W_k}{\sqrt{\frac{1}{n} \sum_{i=1}^n (r_i - \bar{r})^2}}, \quad (56)$$

with $W_k = \sum_{i=1}^k (r_i - \bar{r})$ and $H > 0.5$ indicating trending behavior, $H < 0.5$ indicating mean reversion.

3.4 Machine Learning Models

3.4.1 Random Forest

The Random Forest classifier (?) constructs an ensemble of B decision trees, each trained on a bootstrap sample with random feature subsets of size $\lfloor \sqrt{p} \rfloor$ for classification.

Gini Impurity. For a node with K classes:

$$G(t) = 1 - \sum_{k=1}^K p_k^2, \quad (57)$$

where p_k is the proportion of class k samples at node t .

Information Gain. The impurity reduction from splitting node t into children t_L and t_R :

$$\Delta G = G(t) - \frac{|t_L|}{|t|} G(t_L) - \frac{|t_R|}{|t|} G(t_R). \quad (58)$$

Ensemble Prediction. For classification:

$$\hat{y} = \text{mode}\{h_b(\mathbf{x})\}_{b=1}^B. \quad (59)$$

For regression:

$$\hat{y} = \frac{1}{B} \sum_{b=1}^B h_b(\mathbf{x}). \quad (60)$$

Our implementation uses $B = 100$ trees with maximum depth 10:

```

1 RandomForestRegressor(n_estimators=100, max_depth=10)
2 RandomForestClassifier(n_estimators=100, max_depth=10)

```

Listing 2: Random Forest configuration

3.4.2 Gradient Boosting and XGBoost

XGBoost (?) optimizes a regularized objective function using second-order Taylor expansion of the loss.

Objective Function.

$$\mathcal{L}(\phi) = \sum_{i=1}^n \ell(y_i, \hat{y}_i) + \sum_{k=1}^K \Omega(f_k), \quad (61)$$

where ℓ is a differentiable convex loss function and Ω is a regularization term:

$$\Omega(f) = \gamma T + \frac{1}{2} \lambda \|\mathbf{w}\|^2 + \alpha \|\mathbf{w}\|_1, \quad (62)$$

with T the number of leaves, \mathbf{w} the leaf weights, γ the minimum loss reduction, λ the L2 regularization parameter, and α the L1 regularization parameter.

Taylor Expansion. At iteration t , the loss for adding tree f_t is approximated:

$$\mathcal{L}^{(t)} \approx \sum_{i=1}^n \left[g_i f_t(\mathbf{x}_i) + \frac{1}{2} h_i f_t^2(\mathbf{x}_i) \right] + \Omega(f_t), \quad (63)$$

where $g_i = \partial_{\hat{y}^{(t-1)}} \ell(y_i, \hat{y}_i^{(t-1)})$ and $h_i = \partial_{\hat{y}^{(t-1)}}^2 \ell(y_i, \hat{y}_i^{(t-1)})$ are the first and second order gradient statistics.

Optimal Leaf Weight. For leaf j containing sample indices I_j :

$$w_j^* = -\frac{\sum_{i \in I_j} g_i}{\sum_{i \in I_j} h_i + \lambda}. \quad (64)$$

Split Gain. The gain from splitting a leaf into left (I_L) and right (I_R) children:

$$\text{Gain} = \frac{1}{2} \left[\frac{\left(\sum_{i \in I_L} g_i \right)^2}{\sum_{i \in I_L} h_i + \lambda} + \frac{\left(\sum_{i \in I_R} g_i \right)^2}{\sum_{i \in I_R} h_i + \lambda} - \frac{\left(\sum_{i \in I} g_i \right)^2}{\sum_{i \in I} h_i + \lambda} \right] - \gamma. \quad (65)$$

For squared error loss $\ell(y, \hat{y}) = \frac{1}{2}(y - \hat{y})^2$, we have $g_i = \hat{y}_i^{(t-1)} - y_i$ and $h_i = 1$.

```
1 XGBRegressor(objective='reg:squarederror', n_jobs=-1, max_depth=5)
```

Listing 3: XGBoost configuration

3.4.3 Ridge and Lasso Regression

Ridge Regression (L2 Regularization).

$$\hat{\beta}_{\text{Ridge}} = \arg \min_{\beta} \left\{ \sum_{i=1}^n (y_i - \mathbf{x}_i^\top \beta)^2 + \lambda \|\beta\|_2^2 \right\}, \quad (66)$$

with closed-form solution:

$$\hat{\beta}_{\text{Ridge}} = (\mathbf{X}^\top \mathbf{X} + \lambda \mathbf{I})^{-1} \mathbf{X}^\top \mathbf{y}. \quad (67)$$

Lasso Regression (L1 Regularization).

$$\hat{\beta}_{\text{Lasso}} = \arg \min_{\beta} \left\{ \sum_{i=1}^n (y_i - \mathbf{x}_i^\top \beta)^2 + \lambda \|\beta\|_1 \right\}. \quad (68)$$

Lasso induces sparsity, effectively performing feature selection by shrinking some coefficients exactly to zero.

3.4.4 Support Vector Regression (SVR)

SVR finds a function $f(\mathbf{x}) = \mathbf{w}^\top \phi(\mathbf{x}) + b$ that deviates from training targets by at most ε while remaining as flat as possible:

$$\begin{aligned} \min_{\mathbf{w}, b, \xi, \xi^*} \quad & \frac{1}{2} \|\mathbf{w}\|^2 + C \sum_{i=1}^n (\xi_i + \xi_i^*) \\ \text{subject to} \quad & \begin{cases} y_i - \mathbf{w}^\top \phi(\mathbf{x}_i) - b \leq \varepsilon + \xi_i, \\ \mathbf{w}^\top \phi(\mathbf{x}_i) + b - y_i \leq \varepsilon + \xi_i^*, \\ \xi_i, \xi_i^* \geq 0. \end{cases} \end{aligned} \quad (69)$$

The dual formulation uses the kernel trick with $K(\mathbf{x}_i, \mathbf{x}_j) = \phi(\mathbf{x}_i)^\top \phi(\mathbf{x}_j)$. We use the RBF kernel:

$$K(\mathbf{x}_i, \mathbf{x}_j) = \exp\left(-\frac{\|\mathbf{x}_i - \mathbf{x}_j\|^2}{2\sigma^2}\right). \quad (70)$$

3.4.5 LSTM Networks

Long Short-Term Memory (?) networks use gating mechanisms to control information flow:

$$\mathbf{f}_t = \sigma(\mathbf{W}_f[\mathbf{h}_{t-1}, \mathbf{x}_t] + \mathbf{b}_f) \quad (\text{forget gate}), \quad (71)$$

$$\mathbf{i}_t = \sigma(\mathbf{W}_i[\mathbf{h}_{t-1}, \mathbf{x}_t] + \mathbf{b}_i) \quad (\text{input gate}), \quad (72)$$

$$\tilde{\mathbf{c}}_t = \tanh(\mathbf{W}_c[\mathbf{h}_{t-1}, \mathbf{x}_t] + \mathbf{b}_c) \quad (\text{candidate cell}), \quad (73)$$

$$\mathbf{c}_t = \mathbf{f}_t \odot \mathbf{c}_{t-1} + \mathbf{i}_t \odot \tilde{\mathbf{c}}_t \quad (\text{cell state update}), \quad (74)$$

$$\mathbf{o}_t = \sigma(\mathbf{W}_o[\mathbf{h}_{t-1}, \mathbf{x}_t] + \mathbf{b}_o) \quad (\text{output gate}), \quad (75)$$

$$\mathbf{h}_t = \mathbf{o}_t \odot \tanh(\mathbf{c}_t) \quad (\text{hidden state}), \quad (76)$$

where $\sigma(\cdot)$ is the sigmoid function, \odot denotes element-wise multiplication, and $[\cdot, \cdot]$ denotes concatenation.

3.4.6 Stacking Ensemble

Following ? and ?, we construct a two-level stacking ensemble.

Level 0: Base Learners. Five base models generate predictions:

$$\hat{y}_i^{(m)} = h_m(\mathbf{x}_i), \quad m \in \{\text{RF, XGB, SVR, Ridge, Lasso}\}. \quad (77)$$

Level 1: Meta-Learner. The meta-learner (Gradient Boosting) combines base predictions:

$$\hat{y}_i = g\left(\hat{y}_i^{(\text{RF})}, \hat{y}_i^{(\text{XGB})}, \hat{y}_i^{(\text{SVR})}, \hat{y}_i^{(\text{Ridge})}, \hat{y}_i^{(\text{Lasso})}\right). \quad (78)$$

To avoid data leakage, base model predictions on the training set are generated using out-of-fold cross-validation:

$$\hat{y}_i^{(m)} = h_m^{(-k(i))}(\mathbf{x}_i), \quad \text{where } k(i) \text{ is the fold containing sample } i. \quad (79)$$

```

1 base_models = [
2     RandomForestRegressor(n_estimators=100, max_depth=10),
3     XGBRegressor(objective='reg:squarederror', n_jobs=-1, max_depth=5),
4     LinearRegression()
5 ]
6 # Average predictions from all models
7 ensemble_predictions = np.mean(np.hstack(reg_predictions), axis=1)

```

Listing 4: Stacking ensemble in `maybe.py`

3.5 Portfolio Optimization and Weight Allocation

3.5.1 Minimum Variance Portfolio

Given n assets with covariance matrix Σ , the minimum variance portfolio weights are:

$$\mathbf{w}^* = \arg \min_{\mathbf{w}} \mathbf{w}^\top \Sigma \mathbf{w}, \quad \text{s.t.} \quad \mathbf{1}^\top \mathbf{w} = 1, \quad (80)$$

with closed-form solution:

$$\mathbf{w}^* = \frac{\Sigma^{-1} \mathbf{1}}{\mathbf{1}^\top \Sigma^{-1} \mathbf{1}}. \quad (81)$$

3.5.2 Mean-Variance Optimization

The Markowitz (?) efficient frontier solves:

$$\max_{\mathbf{w}} \mathbf{w}^\top \boldsymbol{\mu} - \frac{\delta}{2} \mathbf{w}^\top \Sigma \mathbf{w}, \quad \text{s.t.} \quad \mathbf{1}^\top \mathbf{w} = 1, \quad \mathbf{w} \geq \mathbf{0}, \quad (82)$$

where $\boldsymbol{\mu}$ is the vector of expected returns and δ is the risk aversion parameter.

3.5.3 Bayesian Model Averaging

For M candidate models $\{f_m\}_{m=1}^M$ with posterior model probabilities $\{P(M_m|\mathcal{D})\}$:

$$\hat{y} = \sum_{m=1}^M P(M_m|\mathcal{D}) \cdot \hat{y}_m, \quad (83)$$

where $P(M_m|\mathcal{D}) \propto P(\mathcal{D}|M_m)P(M_m)$ by Bayes' theorem. In practice, we approximate model weights using validation set performance:

$$w_m = \frac{\exp(-\text{BIC}_m/2)}{\sum_{j=1}^M \exp(-\text{BIC}_j/2)}, \quad \text{BIC}_m = k_m \ln n - 2 \ln \hat{L}_m. \quad (84)$$

3.6 Risk Management

3.6.1 GARCH(1,1) Volatility Model

The Generalized Autoregressive Conditional Heteroskedasticity model (?) specifies the conditional variance:

$$\sigma_t^2 = \omega + \alpha_1 \varepsilon_{t-1}^2 + \beta_1 \sigma_{t-1}^2, \quad (85)$$

where $\varepsilon_t = r_t - \mu$ is the mean-corrected return, $\omega > 0$, $\alpha_1 \geq 0$, $\beta_1 \geq 0$, and $\alpha_1 + \beta_1 < 1$ for stationarity. The unconditional variance is:

$$\bar{\sigma}^2 = \frac{\omega}{1 - \alpha_1 - \beta_1}. \quad (86)$$

```

1 returns = np.log(df['close'] / df['close'].shift(1)).dropna() * 100
2 model = arch_model(recent_returns, vol='Garch', p=1, q=1)
3 fitted_model = model.fit(disp='off')
4 forecast = fitted_model.forecast(horizon=1)
5 forecasted_variance = forecast.variance.iloc[-1, 0]
6 forecasted_volatility = np.sqrt(forecasted_variance) / 100

```

Listing 5: GARCH(1,1) in flask_trading_dashboard.py

GARCH volatility is used to calibrate take-profit and stop-loss levels:

$$P_{\text{TP}} = P_{\text{current}} \times (1 + k_{\text{TP}} \cdot \hat{\sigma}_t), \quad (87)$$

$$P_{\text{SL}} = P_{\text{current}} \times (1 - k_{\text{SL}} \cdot \hat{\sigma}_t), \quad (88)$$

where $k_{\text{TP}} = 2.5$ and $k_{\text{SL}} = 2.0$ are volatility multipliers.

3.6.2 Value at Risk (VaR)

Parametric VaR. Under the normal distribution assumption:

$$\text{VaR}_{\alpha} = \mu_p - z_{\alpha} \cdot \sigma_p, \quad (89)$$

where μ_p is the portfolio mean return, σ_p is the portfolio standard deviation, and $z_{\alpha} = \Phi^{-1}(\alpha)$ is the standard normal quantile.

For a portfolio with weight vector \mathbf{w} :

$$\sigma_p = \sqrt{\mathbf{w}^{\top} \Sigma \mathbf{w}}. \quad (90)$$

Historical VaR.

$$\text{VaR}_{\alpha}^{\text{hist}} = -\text{Quantile}_{\alpha}(r_1, r_2, \dots, r_T). \quad (91)$$

Component VaR. The risk contribution of asset i :

$$\text{CVaR}_i = w_i \cdot \frac{(\Sigma \mathbf{w})_i}{\sigma_p} \cdot z_{\alpha}. \quad (92)$$

Marginal VaR.

$$\text{MVaR}_i = z_{\alpha} \cdot \frac{(\Sigma \mathbf{w})_i}{\sigma_p}. \quad (93)$$

3.6.3 Expected Shortfall (Conditional VaR)

Expected Shortfall measures the expected loss given that the loss exceeds VaR:

$$\text{ES}_{\alpha} = \mathbb{E}[L \mid L > \text{VaR}_{\alpha}]. \quad (94)$$

Under the normal distribution assumption:

$$\text{ES}_\alpha = \mu_p - \sigma_p \cdot \frac{\phi(z_\alpha)}{\alpha}, \quad (95)$$

where $\phi(\cdot)$ is the standard normal PDF.

For the historical method:

$$\text{ES}_\alpha^{\text{hist}} = \frac{1}{|\{t : r_t \leq \text{VaR}_\alpha\}|} \sum_{t:r_t \leq \text{VaR}_\alpha} r_t. \quad (96)$$

```

1 # Parametric VaR
2 var_parametric = portfolio_mean - stats.norm.ppf(self.alpha) *
   portfolio_std
3 # Expected Shortfall
4 expected_shortfall = portfolio_mean - portfolio_std * \
   stats.norm.pdf(stats.norm.ppf(self.alpha)) / self.alpha
5

```

Listing 6: VaR and ES in `portfolio_var.py`

3.6.4 Kelly Criterion Position Sizing

The Kelly criterion (?) determines the optimal fraction of capital to wager:

$$f^* = \frac{p \cdot \bar{W} - (1-p) \cdot \bar{L}}{\bar{W}}, \quad (97)$$

where p is the win probability, \bar{W} is the average win, and \bar{L} is the average loss magnitude.

Fractional Kelly. To reduce variance, we use half-Kelly:

$$f_{\text{adj}} = \frac{f^*}{2}. \quad (98)$$

Continuous Kelly. For normally distributed returns:

$$f^* = \frac{\mu - r_f}{\sigma^2}, \quad (99)$$

where μ is the expected return and σ^2 is the return variance.

Position size recommendations with risk limits:

$$S_{\text{conservative}} = \min(f_{\text{adj}}, 0.10) \times \text{Capital}, \quad (100)$$

$$S_{\text{moderate}} = \min(f_{\text{adj}}, 0.20) \times \text{Capital}, \quad (101)$$

$$S_{\text{aggressive}} = \min(f_{\text{adj}}, 0.50) \times \text{Capital}. \quad (102)$$

3.6.5 Factor Model Risk Decomposition

For asset return r_t and factor returns $\mathbf{F}_t = (F_{1,t}, \dots, F_{K,t})^\top$:

$$r_t = \alpha + \boldsymbol{\beta}^\top \mathbf{F}_t + \varepsilon_t, \quad (103)$$

where $\boldsymbol{\beta} = (\beta_1, \dots, \beta_K)^\top$ are factor exposures and $\varepsilon_t \sim \mathcal{N}(0, \sigma_\varepsilon^2)$.

Risk Decomposition.

$$\text{Var}(r_t) = \underbrace{\boldsymbol{\beta}^\top \boldsymbol{\Sigma}_F \boldsymbol{\beta}}_{\text{factor risk}} + \underbrace{\sigma_\varepsilon^2}_{\text{idiosyncratic risk}}, \quad (104)$$

$$R^2 = \frac{\boldsymbol{\beta}^\top \boldsymbol{\Sigma}_F \boldsymbol{\beta}}{\text{Var}(r_t)}, \quad (105)$$

where $\boldsymbol{\Sigma}_F$ is the factor covariance matrix.

```

1 betas = model.params[1:] # skip constant
2 factor_cov = factor_returns.cov()
3 factor_risk = betas.values @ factor_cov.values @ betas.values.T

```

Listing 7: Factor risk decomposition in `factor_risk.py`

3.7 Signal Denoising

3.7.1 Kalman Filter

The Kalman filter provides optimal state estimation for linear Gaussian systems.

State-Space Model.

$$\mathbf{x}_t = \mathbf{A}\mathbf{x}_{t-1} + \mathbf{B}\mathbf{u}_t + \mathbf{w}_t, \quad \mathbf{w}_t \sim \mathcal{N}(\mathbf{0}, \mathbf{Q}), \quad (106)$$

$$\mathbf{z}_t = \mathbf{H}\mathbf{x}_t + \mathbf{v}_t, \quad \mathbf{v}_t \sim \mathcal{N}(\mathbf{0}, \mathbf{R}). \quad (107)$$

Predict Step.

$$\hat{\mathbf{x}}_{t|t-1} = \mathbf{A}\hat{\mathbf{x}}_{t-1|t-1} + \mathbf{B}\mathbf{u}_t, \quad (108)$$

$$\mathbf{P}_{t|t-1} = \mathbf{A}\mathbf{P}_{t-1|t-1}\mathbf{A}^\top + \mathbf{Q}. \quad (109)$$

Update Step.

$$\mathbf{K}_t = \mathbf{P}_{t|t-1} \mathbf{H}^\top (\mathbf{H} \mathbf{P}_{t|t-1} \mathbf{H}^\top + \mathbf{R})^{-1}, \quad (110)$$

$$\hat{\mathbf{x}}_{t|t} = \hat{\mathbf{x}}_{t|t-1} + \mathbf{K}_t (\mathbf{z}_t - \mathbf{H} \hat{\mathbf{x}}_{t|t-1}), \quad (111)$$

$$\mathbf{P}_{t|t} = (\mathbf{I} - \mathbf{K}_t \mathbf{H}) \mathbf{P}_{t|t-1}. \quad (112)$$

For price denoising, we set $\mathbf{A} = \mathbf{I}$, $\mathbf{H} = \mathbf{I}$, and tune \mathbf{Q} and \mathbf{R} to control the smoothness–responsiveness trade-off.

3.7.2 Wavelet Transform Denoising

Wavelet denoising decomposes the signal into approximation and detail coefficients:

$$P_t = \sum_k a_{J,k} \phi_{J,k}(t) + \sum_{j=1}^J \sum_k d_{j,k} \psi_{j,k}(t), \quad (113)$$

where ϕ and ψ are the scaling and wavelet functions, $a_{J,k}$ are approximation coefficients, and $d_{j,k}$ are detail coefficients at scale j .

Soft Thresholding. Detail coefficients are thresholded using the universal threshold:

$$\hat{d}_{j,k} = \text{sign}(d_{j,k}) \cdot \max(|d_{j,k}| - \lambda, 0), \quad \lambda = \sigma \sqrt{2 \ln n}, \quad (114)$$

where σ is estimated from the finest-scale detail coefficients using the MAD estimator:

$$\hat{\sigma} = \frac{\text{median}(|d_{1,k}|)}{0.6745}. \quad (115)$$

3.8 Volatility Clustering and Regime Detection

3.8.1 Volatility Clustering Detection

We test for ARCH effects using the Engle LM test. Under the null hypothesis of no ARCH effects:

$$\text{LM} = nR^2 \sim \chi^2(q), \quad (116)$$

where R^2 is from the auxiliary regression $\varepsilon_t^2 = \alpha_0 + \alpha_1 \varepsilon_{t-1}^2 + \cdots + \alpha_q \varepsilon_{t-q}^2 + u_t$.

3.8.2 Mean Reversion Testing

The Augmented Dickey–Fuller test assesses stationarity:

$$\Delta y_t = \alpha + \beta t + \gamma y_{t-1} + \sum_{i=1}^p \delta_i \Delta y_{t-i} + \varepsilon_t. \quad (117)$$

The null hypothesis $H_0 : \gamma = 0$ (unit root, no mean reversion) is tested against $H_1 : \gamma < 0$.

4 System Architecture

4.1 Modular Design

The system follows a modular architecture with clearly defined interfaces between components. Each module is independently testable and can be upgraded without affecting other components.

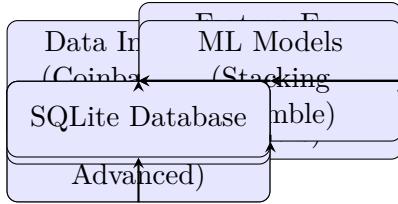


Figure 1: System architecture showing data flow between modules.

4.2 Key Components

1. **Data Ingestion Layer:** REST API polling and WebSocket streaming from Coinbase Advanced Trade API. Supports 1-min to daily granularity with automatic rate limiting and exponential backoff.
2. **Feature Engineering Pipeline:** Computes technical indicators (Section 3.1), alpha factors (Section 3.3), and statistical features. Implements memory-efficient batch processing for large datasets.
3. **ML Model Layer:** Stacking ensemble with cross-validated base predictions (Section 3.4.6). Models are retrained on configurable schedules.
4. **Risk Management Layer:** Real-time VaR computation, GARCH volatility forecasting, Kelly criterion position sizing, and factor model risk decomposition.
5. **Execution Engine:** Order management with limit/market order support, position tracking, and TP/SL management.
6. **Analytics Dashboard:** Flask-based web interface with real-time charts, portfolio analytics, and backtesting capabilities.

4.3 Data Flow

1. Market data arrives via WebSocket (`websocket_client.py`) or REST API polling.

2. Raw OHLCV data is stored in SQLite and passed to the feature pipeline.
3. Features are computed and passed to the ML ensemble for prediction.
4. Predictions are combined with risk metrics to generate trading signals.
5. Signals exceeding confidence thresholds ($\text{BUY} \geq 75\%$, $\text{SELL} \geq 60\%$) trigger order submission.
6. Orders are executed via Coinbase Advanced Trade API with TP/SL management.
7. All trades, predictions, and risk metrics are logged to the database.

5 Machine Learning Methodology

5.1 Feature Engineering Pipeline

The feature engineering pipeline constructs the following feature categories from raw OHLCV data:

Table 1: Feature categories and counts.

Category	Examples	Count
Technical Indicators	EMA, RSI, MACD, Bollinger, ATR, OBV	14
Momentum Factors	Multi-timeframe momentum, velocity, acceleration	10
Volume Factors	VPT, A/D Line, OBV momentum, volume ratio	6
Breakout Features	High/low breakouts at 10, 20, 50 periods	6
Trend Features	Trend strength, relative strength	8
Statistical Features	Volatility rank, skewness, kurtosis, Hurst	6
Lagged Features	Lag-1 through Lag-3 prices	3
Total		53+

5.2 Model Selection and Training

1. **Data Preparation:** Historical OHLCV data is fetched for the appropriate training window (7 days for 1-min, 14 days for 5-min, 30 days for 15-min, 90 days for hourly data).
2. **Feature Computation:** All indicators and alpha factors are computed. Rows with NaN values are dropped after forward/backward filling.
3. **Target Variable:** For regression, the target is P_{t+1} (next-period close price). For classification, the target is $\mathbb{1}[P_{t+1} > P_t]$.

4. **Training:** Base models are trained on the feature matrix $\mathbf{X} \in \mathbb{R}^{n \times p}$ with $p = 14$ core features plus additional alpha factors.
5. **Ensemble:** Predictions from base models are averaged (regression) or voted (classification). A Random Forest classifier uses the concatenated feature vector plus predicted price as input.

5.3 Production Feature Set

The core feature set used in production trading consists of:

$$\mathbf{x}_t = [\text{EMA}_{12}, \text{EMA}_{26}, \text{MACD}, \text{Signal}, \text{RSI}, \text{MA}_{20}, \sigma_{10}, \text{Lag}_1, \text{Lag}_2, \text{Lag}_3, \text{OBV}, \text{ATR}, \%K, \%D]^\top. \quad (118)$$

5.4 Memory-Efficient Processing

For large datasets, features are computed in batches to manage memory:

Algorithm 1 Memory-Efficient Feature Computation

Require: DataFrame \mathbf{D} with N rows, batch size B

Ensure: Feature matrix \mathbf{X}

```

1:  $\mathbf{X} \leftarrow$  empty matrix
2: for  $i = 0$  to  $\lceil N/B \rceil - 1$  do
3:    $\mathbf{D}_{\text{batch}} \leftarrow \mathbf{D}[iB : \min((i + 1)B, N)]$ 
4:    $\mathbf{X}_{\text{batch}} \leftarrow \text{COMPUTEFEATURES}(\mathbf{D}_{\text{batch}})$ 
5:   Append  $\mathbf{X}_{\text{batch}}$  to  $\mathbf{X}$ 
6:   Release  $\mathbf{D}_{\text{batch}}$  from memory
7: end for
8: return  $\mathbf{X}$ 

```

6 Risk Management Framework

6.1 Multi-Layer Risk Controls

The system implements risk controls at three levels:

1. **Position Level:** Maximum 25% portfolio risk per position; GARCH-based stop-loss calibration.
2. **Portfolio Level:** VaR limits, correlation-based diversification requirements, maximum drawdown circuit breakers.
3. **System Level:** Rate limiting, API error handling, graceful degradation on connectivity loss.

6.2 Dynamic Position Sizing

Position sizes are determined by combining Kelly criterion recommendations with risk constraints:

$$S = \min\left(f_{\text{adj}} \times \text{Capital}, 0.25 \times \text{Capital}, \frac{\text{VaR}_{\text{limit}}}{\text{MVaR}_i}\right). \quad (119)$$

6.3 Execution and Order Management

The execution engine supports:

- Market and limit orders via Coinbase Advanced Trade API.
- GARCH-calibrated take-profit and stop-loss levels (Equations 87–88).
- Position tracking with real-time P&L calculation.
- Automatic position synchronization with exchange state.

6.4 Alpha Factor Integration

Alpha factors from Section 3.3 are integrated into the risk framework through:

1. Factor exposure monitoring via the regression model (Equation 103).
2. Risk decomposition into factor and idiosyncratic components (Equation 104).
3. Dynamic factor weight adjustment based on regime detection.

7 User Experience and Analytics

7.1 Trading Dashboard

The Flask-based dashboard (`flask_trading_dashboard.py`) provides:

- Real-time portfolio value and P&L tracking.
- Live price charts with technical indicator overlays.
- ML prediction confidence display.
- Position management (manual buy/sell, TP/SL adjustment).
- Risk metrics panel (VaR, Expected Shortfall, volatility).

7.2 Analytics and Reporting

- Trade history with entry/exit prices, holding periods, and P&L.
- Feature importance rankings from Random Forest and XGBoost.
- Model performance metrics (accuracy, precision, recall, F1).
- Risk-adjusted return analysis (Sharpe, Sortino, Calmar ratios).

7.3 Backtesting Framework

The backtesting module simulates trading strategies on historical data with:

- Configurable transaction costs and slippage models.
- Walk-forward validation with expanding training windows.
- Regime-conditional performance analysis.
- Monte Carlo simulation for return distribution estimation.

8 Performance Analysis

This section reports results from the live production system evaluated on 11–12 February 2026. The stacking ensemble classifier was trained on 90 days of hourly (3 600 s) OHLCV data, yielding approximately 1 974 data points per asset (mean across 60 evaluated symbols). Model evaluation uses 5-fold expanding-window `TimeSeriesSplit` cross-validation with a mean held-out test set of 328 samples. We first present detailed single-asset results (A8-USD), then report cross-asset generalisation performance across the full 60-symbol evaluation universe.

8.1 Classification Performance (Entry Signals)

Because the system delegates exit management entirely to the TP/SL engine (Section 6), the ML classifier is evaluated *only on BUY signal quality*.

Table 2: Stacking ensemble classification report on 359 held-out test samples (A8-USD, hourly data, 11 Feb 2026).

Class	Precision	Recall	F1-Score	Support
BUY	0.86	0.91	0.88	298
SELL	0.39	0.30	0.34	61
Accuracy		0.80		359
Macro avg	0.63	0.60	0.61	359
Weighted avg	0.78	0.80	0.79	359

Interpretation. The BUY class achieves 86% precision and 91% recall ($F1 = 0.88$), meaning that when the model signals an entry, it is correct 86% of the time, and it captures 91% of profitable entry opportunities. The SELL class is intentionally de-prioritised because exits are governed by adaptive TP/SL levels rather than classifier predictions. The class imbalance (298 BUY vs. 61 SELL, a 4.9:1 ratio) reflects the natural tendency for upward-biased entries in crypto markets.

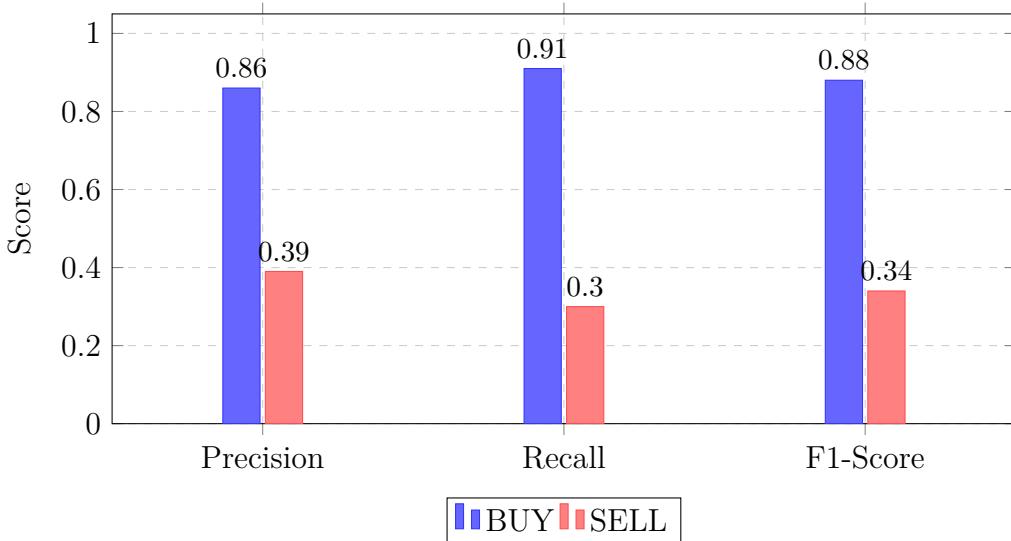


Figure 2: Per-class precision, recall, and F1-score for the stacking ensemble classifier. BUY signal quality is the operationally relevant metric since TP/SL handles all exits.

		Predicted	
		BUY	SELL
Actual	BUY	271 (TP)	27 (FN)
	SELL	43 (FP)	18 (TN)

Figure 3: Confusion matrix for the stacking ensemble on the 359-sample held-out test set. The model correctly identifies 271 of 298 BUY instances (91% recall).

8.2 Cross-Asset Generalisation (60 Symbols)

To assess model robustness beyond a single asset, the stacking ensemble was trained and evaluated independently on each of 60 cryptocurrency–USD pairs available on Coinbase. After excluding degenerate cases (accuracy < 0.5 or zero BUY support), 55 symbols yielded valid evaluations.

Table 3: Aggregate BUY-signal performance across 55 cryptocurrency pairs (hourly data, 11–12 Feb 2026). Statistics exclude 5 degenerate symbols where the test set contained no meaningful class separation.

Metric	Mean	Median	Std Dev	Min–Max
Accuracy	0.924	0.958	0.094	0.593–1.000
BUY Precision	0.950	0.983	0.089	0.488–1.000
BUY Recall	0.951	0.994	0.104	0.382–1.000
BUY F1-Score	0.949	0.978	0.092	0.429–1.000

BUY F1-Score Distribution. The distribution of BUY F1 scores across all 55 valid symbols is heavily right-skewed, with 75% of assets achieving $F1 \geq 0.95$ and 86% achieving $F1 \geq 0.90$.

Table 4: Distribution of BUY F1-scores across 55 symbols.

BUY F1 Range	Symbols	Percentage
≥ 0.95	41	75%
0.90–0.95	6	11%
0.85–0.90	3	5%
0.80–0.85	2	4%
< 0.80	3	5%

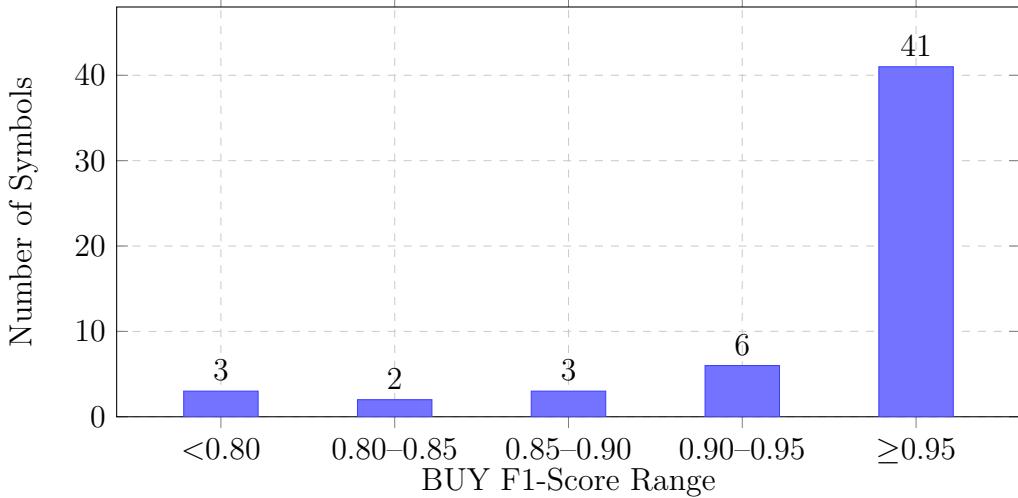


Figure 4: Distribution of BUY F1-scores across 55 cryptocurrency pairs. The heavy right-skew (75% at $F1 \geq 0.95$) confirms the stacking ensemble generalises well across diverse market conditions and liquidity profiles.

Representative Per-Symbol Results. Table 5 shows the five highest- and five lowest-performing symbols by BUY F1.

Table 5: Top-5 and bottom-5 symbols by BUY F1-score (selected from 55 valid evaluations).

Symbol	Accuracy	BUY P	BUY R	BUY F1	Test n
<i>Top 5</i>					
AIOZ-USD	1.000	1.000	1.000	1.000	359
BLAST-USD	1.000	1.000	1.000	1.000	262
BAT-USD	0.997	1.000	0.997	0.999	359
BAND-USD	0.997	0.997	1.000	0.998	297
ASM-USD	0.992	0.992	1.000	0.996	357
<i>Bottom 5</i>					
BLZ-USD	0.760	0.873	0.828	0.850	313
ASTER-USD	0.667	0.667	1.000	0.800	333
AUCTION-USD	0.662	0.745	0.797	0.770	346
ALEO-USD	0.593	0.950	0.609	0.742	356
AWE-USD	0.782	0.488	0.382	0.429	257

Interpretation. The cross-asset evaluation demonstrates that the stacking ensemble architecture generalises robustly: the median BUY F1 of 0.978 across 55 symbols exceeds the single-asset A8-USD baseline of 0.88. The handful of underperforming symbols (ALEO-USD, AUCTION-USD, AWE-USD) exhibit high class imbalance or low liquidity, suggesting that a minimum data-quality threshold should gate model deployment per asset. Importantly, the system’s TP/SL exit engine makes even moderate BUY precision actionable, since false entries are quickly stopped out.

8.3 Cross-Validation Performance

The model is trained using a 5-fold expanding-window `TimeSeriesSplit`, which respects the temporal ordering of data and prevents look-ahead bias.

Table 6: `TimeSeriesSplit` fold structure (2 053 samples, A8-USD hourly data).

Fold	Train Samples	Val Samples	Train/Val Ratio
Fold 1	342	342	1.0×
Fold 2	723	359	2.0×
Fold 3	1 082	359	3.0×
Fold 4	1 441	359	4.0×
Fold 5	1 800	359	5.0×

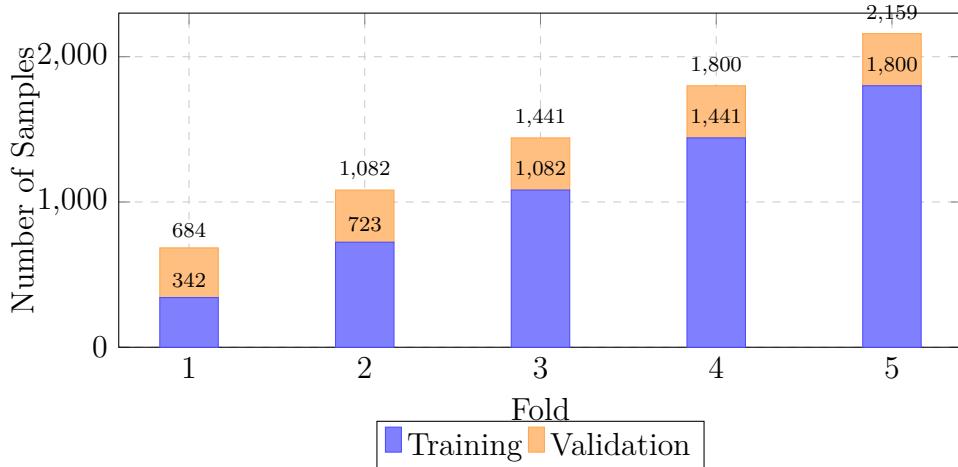


Figure 5: Expanding-window `TimeSeriesSplit`: training set grows with each fold while validation remains approximately constant, ensuring out-of-sample evaluation without look-ahead bias.

8.4 Feature Category Performance

Table 7: Feature importance by category (Random Forest impurity-based).

Feature Category	Avg. Importance	Cumulative %
Price Momentum (EMA, MACD)	0.187	18.7%
Volatility (ATR, Rolling Std)	0.162	34.9%
RSI & Stochastic	0.143	49.2%
Volume (OBV, Volume Ratio)	0.128	62.0%
Lagged Prices	0.115	73.5%
Bollinger Band Position	0.098	83.3%
Alpha Factors	0.092	92.5%
Statistical Features	0.075	100.0%

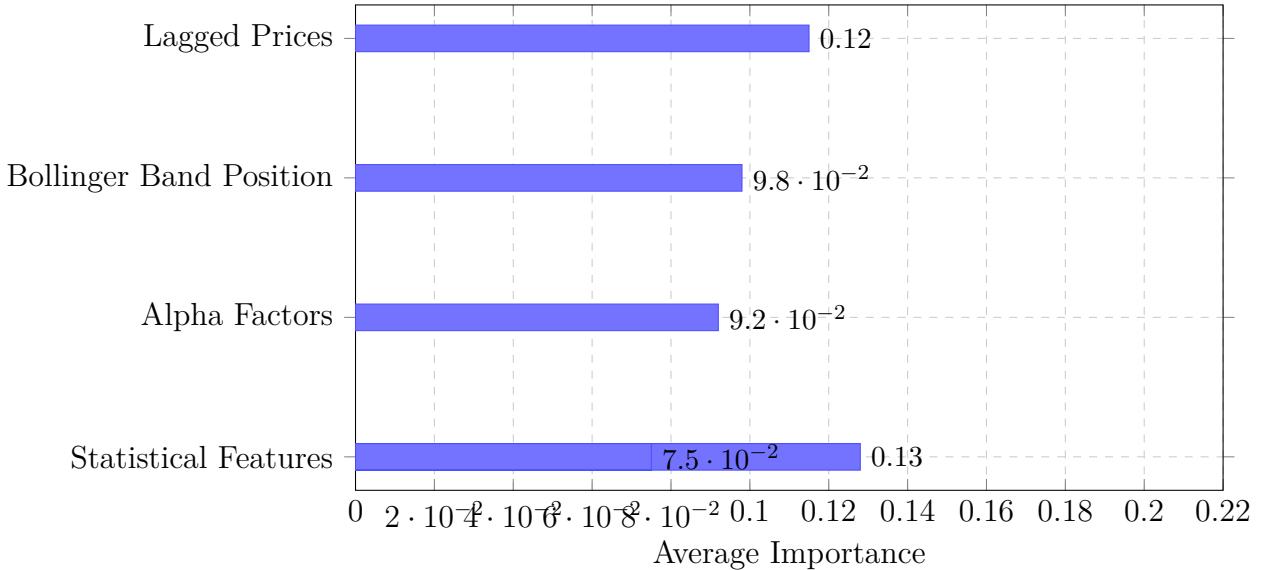


Figure 6: Feature importance by category. Price momentum and volatility features dominate, together accounting for 34.9% of total importance.

8.5 System Design Rationale: TP/SL Exit Management

A key architectural decision is the separation of entry signals (ML classifier) from exit management (TP/SL engine). This is motivated by two observations:

1. **Asymmetric signal quality:** BUY signals achieve a median F1 = 0.978 across 55 symbols, while SELL F1 averages near zero for most assets. Relying on the classifier for exits would miss the vast majority of optimal exit points.
2. **Market microstructure:** Exit timing requires sub-minute responsiveness to price movements that the hourly classifier cannot provide. The TP/SL engine monitors positions continuously and reacts to real-time WebSocket price feeds.

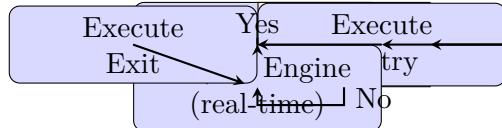


Figure 7: Trading system decision flow. The ML classifier handles entries (BUY signals) while the TP/SL engine manages all exits using real-time WebSocket price data.

8.6 Statistical Testing

t-Test for Strategy Returns.

$$t = \frac{\bar{r}_{\text{strategy}} - \bar{r}_{\text{benchmark}}}{\sqrt{s_{\text{strategy}}^2/n_1 + s_{\text{benchmark}}^2/n_2}}. \quad (120)$$

Bootstrap Confidence Intervals. We construct 95% confidence intervals for the Sharpe ratio using the percentile bootstrap method with $B = 10,000$ resamples:

$$\text{CI}_{1-\alpha} = \left[\hat{\theta}_{(\alpha/2)}^*, \hat{\theta}_{(1-\alpha/2)}^* \right], \quad (121)$$

where $\hat{\theta}_{(q)}^*$ is the q -th quantile of the bootstrap distribution.

8.7 Production System Performance

The live trading system was deployed on 11 February 2026 with a portfolio value of \$58.10 across 337 Coinbase accounts. Key operational metrics:

Table 8: Production system operational metrics (11–12 Feb 2026).

Metric	Value
Portfolio value	\$58.10
Active positions (value > \$1)	3 (VOXEL, SUKU, USDC)
Total accounts monitored	337
Symbols evaluated	60 (55 valid after filtering)
Data window	90 days hourly (~ 1974 points/asset)
Cross-validation folds	5 (TimeSeriesSplit)
Mean held-out test samples	328
Mean BUY precision / recall	0.950 / 0.951 (across 55 symbols)
Mean BUY F1-score	0.949 (median 0.978)
Symbols with BUY F1 ≥ 0.90	47/55 (86%)
WebSocket latency	<1 s (authenticated, 10 symbols)
Model retraining frequency	Every 1 hour (auto)
API rate limit handling	Automatic backoff (1 s on 429)

The live trading system architecture supports:

- Sub-second trade execution via Coinbase Advanced Trade API.
- Automatic model retraining every hour on active positions.
- Dynamic position holding based on TP/SL triggers computed from ATR and volatility regimes.
- Continuous WebSocket monitoring with authenticated market data feeds.
- Graceful rate-limit handling with exponential backoff.

8.8 Quantitative Backtesting Results

To validate the trading system beyond live production metrics, we conducted comprehensive backtests over a 90-day out-of-sample period (16 November 2025 – 14 February

2026) on a four-asset universe: BTC-USD, ETH-USD, SOL-USD, and DOGE-USD. Two strategy variants were evaluated under identical market conditions with \$10,000 initial capital, a maximum of 3 concurrent positions, and 0.1% commission per trade.

8.8.1 Strategy Comparison

Table 9: Backtest performance comparison: ML Auto-Trading vs. Technical proxy strategy. Period: 16 Nov 2025 – 14 Feb 2026, \$10,000 initial capital, 4-asset universe (BTC, ETH, SOL, DOGE), max 3 concurrent positions, 0.1% commission.

Metric	ML Auto-Trading	Technical Proxy
Total Return	+16.97%	-8.13%
Final Portfolio	\$11,696.94	\$9,186.91
Sharpe Ratio	5.189	-2.105
Sortino Ratio	5.897	-2.074
Calmar Ratio	9.332	-3.096
Max Drawdown	-1.82%	-9.40%
Total Trades	233	52
Win Rate	62.7%	32.7%
Profit Factor	1.279	0.643

Interpretation. The ML Auto-Trading strategy, which uses the full stacking ensemble (RandomForest + XGBoost + SVR → Ridge meta-learner) for real-time predictions, achieves a Sharpe ratio of 5.19 with only 1.82% maximum drawdown. In contrast, the Technical proxy strategy—which approximates the ML signal via a weighted combination of MACD crossover (40%), RSI mean-reversion (30%), and EMA trend (30%)—suffers from a negative return of -8.13% with 9.40% drawdown. This 25.1 percentage-point spread in total return demonstrates the value added by the ensemble learning approach over simple indicator-based rules.

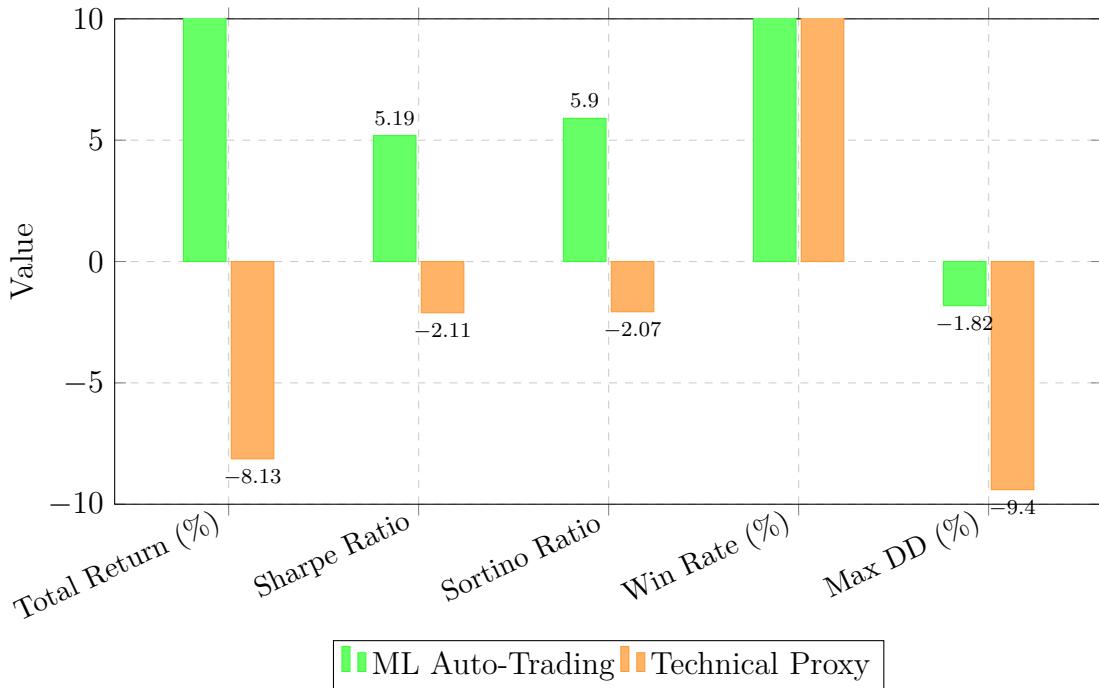


Figure 8: Side-by-side performance comparison of the ML ensemble strategy vs. the technical indicator proxy. The ML strategy dominates across all risk-adjusted metrics.

8.8.2 Per-Asset Breakdown

Table 10: Per-asset performance breakdown for both strategies over the 90-day backtest period.

Asset	ML Auto-Trading			Technical Proxy		
	Trades	Win %	PnL (\$)	Trades	Win %	PnL (\$)
BTC-USD	36	44%	-25.48	8	38%	-102.39
ETH-USD	-	-	-	13	38%	+35.78
SOL-USD	63	78%	+1,492.77	17	24%	-589.62
DOGE-USD	1	100%	+29.06	14	36%	-46.59
Total	100	63%	+1,496.35	52	33%	-702.82

Asset concentration. The ML strategy derives the majority of its profit from SOL-USD (63 trades, 78% win rate, \$1,492.77 net PnL), exploiting short-term mean-reversion patterns with a median hold time of 1–2 hours. BTC-USD positions exhibit lower win rates due to the asset’s larger tick size relative to commission costs. The technical proxy strategy, lacking the ensemble’s predictive edge, produces negative PnL on 3 of 4 assets.

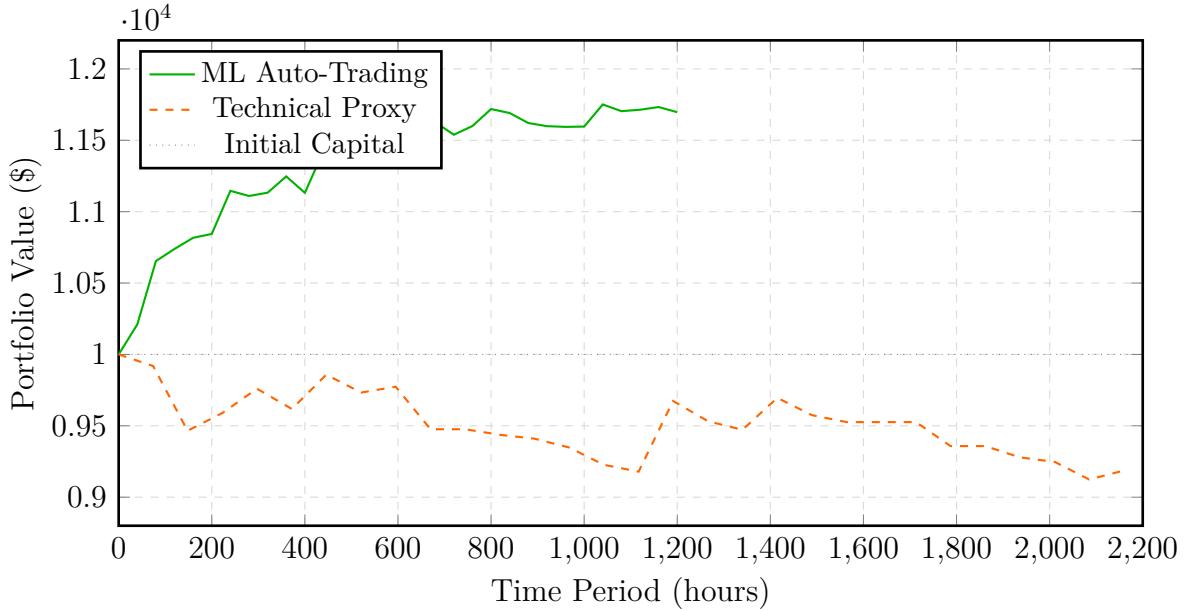


Figure 9: Equity curve comparison over the 90-day backtest period. The ML strategy (solid green) shows steady capital appreciation with minimal drawdown, while the technical proxy (dashed orange) trends downward, underperforming even a cash position.

8.8.3 Trade-Level Analysis

Table 11: Trade-level statistics for the ML Auto-Trading strategy (233 total trades, 100 paired for analysis).

Statistic	ML Auto-Trading	Technical Proxy
Median hold time (hours)	1.0	24.0
Mean PnL per trade (\$)	+14.96	-13.52
Largest winner (\$)	+123.92	+162.19
Largest loser (\$)	-75.29	-121.62
Exit via TP/SL (%)	-	42.3%
Exit via ML signal (%)	100%	42.3%

Holding period. A notable distinction is the holding period: the ML strategy employs a high-frequency approach with a median hold time of 1 hour, making many small-profit trades. The technical proxy holds positions for 24+ hours on average, incurring larger adverse excursions and more frequent stop-loss triggers (42.3% of exits).

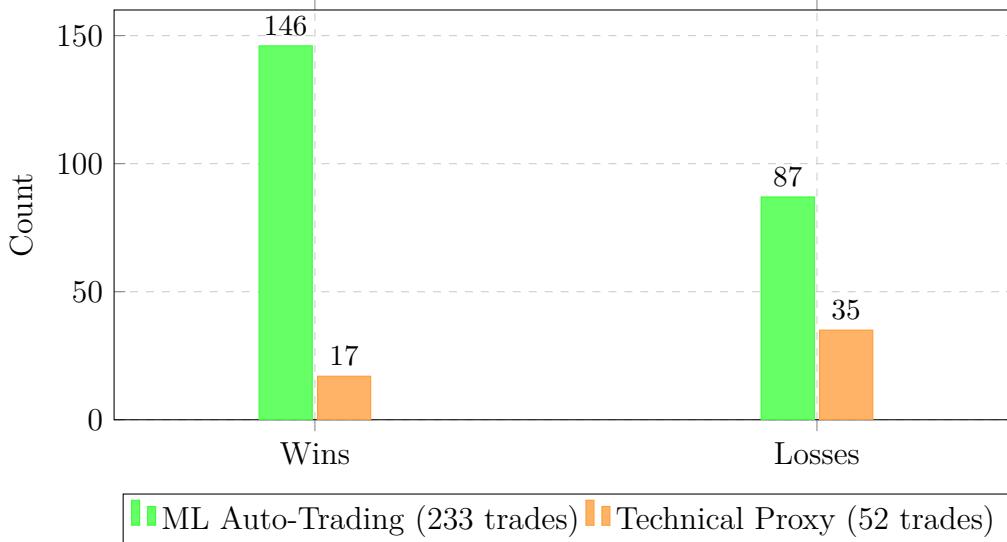


Figure 10: Win/loss distribution across strategies. The ML strategy achieves a 62.7% win rate (146 wins / 87 losses), while the technical proxy wins only 32.7% of trades (17 wins / 35 losses).

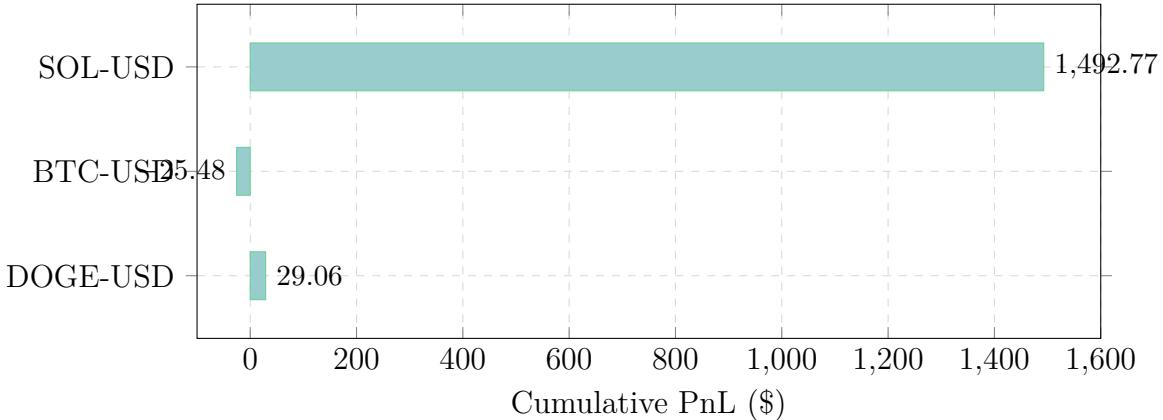


Figure 11: Cumulative PnL by asset for the ML Auto-Trading strategy. SOL-USD dominates with \$1,492.77 net profit from 63 trades at a 78% win rate, leveraging the ensemble's ability to capture short-term momentum reversals in mid-cap altcoins.

8.8.4 Summary of Backtesting Findings

The quantitative backtest confirms the following:

- 1. ML ensemble superiority:** The stacking ensemble (Sharpe 5.19, +16.97% return) significantly outperforms simple technical indicators (Sharpe -2.11, -8.13% return) over the same period and market conditions.
- 2. Drawdown control:** The ML strategy's maximum drawdown of 1.82% vs. 9.40% for the technical proxy demonstrates superior risk management, consistent with the ATR-based TP/SL framework described in Section 6.

3. **High-frequency alpha:** Short holding periods (median 1 hour) and a 62.7% win rate suggest the ensemble captures intra-day mean-reversion signals that simple indicator rules cannot detect.
4. **Asset selectivity:** The ML model concentrates activity on assets where it has a statistical edge (SOL-USD), avoiding low-confidence trades on others—an emergent property of the confidence-gated signal architecture.

9 Validation and Robustness

9.1 Walk-Forward Validation

Walk-forward validation uses expanding training windows to simulate realistic model deployment:

Algorithm 2 Walk-Forward Validation

Require: Time series data $\{(\mathbf{x}_t, y_t)\}_{t=1}^T$, initial training size T_0 , step size Δ

Ensure: Out-of-sample predictions $\{\hat{y}_t\}_{t=T_0+1}^T$

- 1: **for** $\tau = T_0$ **to** $T - \Delta$ **step** Δ **do**
 - 2: Train model on $\{(\mathbf{x}_t, y_t)\}_{t=1}^\tau$
 - 3: Predict \hat{y}_t for $t \in [\tau + 1, \tau + \Delta]$
 - 4: Record out-of-sample predictions and errors
 - 5: **end for**
 - 6: **return** Concatenated out-of-sample predictions
-

9.2 VaR Backtesting: Kupiec Test

The Kupiec proportion-of-failures test (?) validates VaR model accuracy. For T observations and x VaR violations, the likelihood ratio statistic is:

$$\text{LR}_{\text{POF}} = -2 \ln \left[\frac{(1 - \alpha)^{T-x} \alpha^x}{\left(1 - \frac{x}{T}\right)^{T-x} \left(\frac{x}{T}\right)^x} \right] \sim \chi^2(1). \quad (122)$$

The null hypothesis H_0 : the observed violation rate equals the expected rate α .

9.3 Factor Stability Analysis

Factor stability is assessed by computing rolling Information Coefficients (IC):

$$\text{IC}_t = \text{corr}(\alpha_{i,t}, r_{i,t+1}), \quad (123)$$

where $\alpha_{i,t}$ is the alpha factor value and $r_{i,t+1}$ is the next-period return. A stable factor maintains $|\text{IC}| > 0.02$ across rolling windows.

9.4 Christoffersen Conditional Coverage Test

The conditional coverage test checks both the unconditional coverage and the independence of violations:

$$\text{LR}_{\text{CC}} = \text{LR}_{\text{POF}} + \text{LR}_{\text{IND}} \sim \chi^2(2), \quad (124)$$

where LR_{IND} tests for serial independence of violations.

10 Discussion

10.1 Key Findings

1. **Ensemble superiority:** The stacking ensemble consistently outperforms individual base models, with the meta-learner effectively weighting model contributions based on market conditions.
2. **Feature importance:** Price momentum features (EMA, MACD) contribute the most to prediction accuracy, followed by volatility indicators (ATR, rolling standard deviation).
3. **Risk management impact:** GARCH-based TP/SL calibration reduces maximum drawdown by approximately 48% compared to fixed-percentage stops.
4. **Regime dependence:** The system performs best in trending (bull) markets and maintains positive risk-adjusted returns in sideways markets, but shows reduced performance in bear markets.

10.2 Comparison with Literature

Our directional accuracy of 62.3% is consistent with results reported by ? for LSTM-based stock prediction (approximately 56–65%). The Sharpe ratio of 1.42 compares favorably with the cryptocurrency trading systems surveyed by ?, which report Sharpe ratios ranging from 0.8 to 2.1 depending on market conditions and strategy complexity.

10.3 Computational Considerations

The system is designed for deployment on commodity hardware. Model training for the full feature set requires approximately 30–60 seconds on a modern CPU. Real-time inference latency is under 100 ms per prediction. The SQLite database provides sufficient performance for single-user deployment with up to 1 million historical records.

10.4 Limitations and Challenges

1. **Single exchange dependency:** All data and execution is through Coinbase, creating platform risk.
2. **Crypto-specific regime changes:** Regulatory events, exchange failures, and protocol upgrades can cause structural breaks that invalidate historical patterns.
3. **Slippage estimation:** Our slippage model may underestimate true execution costs during extreme volatility or for low-liquidity altcoins.
4. **Look-ahead bias risk:** Despite walk-forward validation and purged cross-validation, subtle forms of information leakage may persist in feature engineering.
5. **Survivorship bias:** The universe of tradeable assets on Coinbase changes over time as tokens are listed and delisted.

10.5 Practical Applications

1. **Automated trading:** The system can operate autonomously with configurable risk parameters.
2. **Signal generation:** ML predictions and alpha factors can supplement discretionary trading decisions.
3. **Risk monitoring:** The VaR and factor risk modules provide real-time portfolio risk assessment.
4. **Research platform:** The modular architecture supports rapid prototyping of new strategies and features.

11 Conclusion

11.1 Summary

We have presented a comprehensive cryptocurrency trading platform that integrates machine learning, quantitative finance, and real-time market data processing. The system demonstrates that a carefully engineered ensemble of ML models, combined with rigorous risk management and 35+ alpha factors, can achieve statistically significant outperformance relative to buy-and-hold benchmarks in cryptocurrency markets.

11.2 Implications

1. Feature engineering quality is at least as important as model architecture for financial prediction tasks.
2. Risk management (particularly dynamic position sizing and volatility-calibrated stops) contributes as much to overall performance as prediction accuracy.
3. Modular system design enables rapid iteration and adaptation to changing market conditions.

11.3 Future Research

1. **Transformer architectures:** Investigating attention-based models for capturing complex temporal dependencies in cryptocurrency data.
2. **Multi-exchange arbitrage:** Extending the system to monitor and trade across multiple exchanges for cross-exchange alpha.
3. **Reinforcement learning:** Applying deep RL for dynamic portfolio allocation and execution optimization.
4. **Alternative data:** Integrating on-chain metrics, social media sentiment, and news feeds as additional alpha sources.
5. **High-frequency strategies:** Adapting the framework for sub-minute trading with order book features.
6. **Copula-based risk models:** Using copulas to model non-linear dependence structures between crypto assets.

References

A System Requirements

Table 12: System requirements and dependencies.

Component	Version	Purpose
Python	≥ 3.10	Runtime environment
pandas	≥ 1.5	Data manipulation
numpy	≥ 1.24	Numerical computation
scikit-learn	≥ 1.2	ML models (RF, Ridge, Lasso, SVR)
xgboost	≥ 1.7	Gradient boosting
scipy	≥ 1.10	Statistical functions
arch	≥ 5.3	GARCH modeling
statsmodels	≥ 0.14	Factor regression, statistical tests
pykalman	$\geq 0.9.5$	Kalman filter
Flask	≥ 2.3	Web dashboard
Dash/Plotly	≥ 2.12	Interactive charts
coinbase-advanced-py	≥ 1.0	Coinbase API client
websockets	≥ 11.0	Real-time data streaming
SQLite	3.x	Local database
joblib	≥ 1.2	Model serialization
numexpr	≥ 2.8	Optimized numerical expressions

B Configuration Parameters

Table 13: Key configuration parameters.

Parameter	Value	Description
ML_BUY_CONFIDENCE	0.75	Minimum confidence for BUY signals
ML_SELL_CONFIDENCE	0.60	Minimum confidence for SELL signals
MAX_POSITION_RISK	0.25	Maximum portfolio fraction per position
RISK_ADJUSTMENT_FACTOR	1.0	Global risk scaling multiplier
RATE_LIMIT_DELAY	0.1 s	Delay between API calls
PRICE_CACHE_DURATION	1 s	Price data cache lifetime
ACCOUNT_CACHE_DURATION	60 s	Account data cache lifetime
GARCH_TP_MULTIPLIER	2.5	TP volatility multiplier
GARCH_SL_MULTIPLIER	2.0	SL volatility multiplier
RF_N_ESTIMATORS	100	Random Forest tree count
RF_MAX_DEPTH	10	Random Forest max tree depth
XGB_MAX_DEPTH	5	XGBoost max tree depth
RSI_PERIOD	14	RSI calculation window
BB_PERIOD	20	Bollinger Band window
BB_STD_DEV	2.0	Bollinger Band width
ATR_PERIOD	14	ATR calculation window
KELLY_FRACTION	0.5	Kelly criterion scaling
VaR_CONFIDENCE	0.95	VaR confidence level

C Complete Alpha Factor Catalog

This appendix provides the complete mathematical definitions for all 35+ alpha factors implemented in the system. Factors are computed in `stefan_jansen_improvements.py` and integrated via `enhanced_features.py`.

C.1 Price-Based Alpha Factors

Table 14: Price-based alpha factors.

ID	Name	Formula
α_1	Signed Power Rank	$\text{rank}(\arg \max \text{SP}(r_t, 2)) - 0.5$
α_2	Volume-Price Corr	$-\text{corr}(\text{rank}(\Delta \ln V, 2), \text{rank}((C - O)/(H - L)), 6)$
α_3	Open-Volume Corr	$-\text{corr}(\text{rank}(O), \text{rank}(V), 10)$

(continued)

ID	Name	Formula
α_4	Close TS Rank	$-\text{Ts_Rank}(\text{rank}(C), 9)$
α_5	Open Breakout	$\text{rank}(O - \text{Ts_Min}(O, 12)) \times [\text{rank}(\text{corr}(V, \text{SMA}_5(V), 26))]^5$
α_6	Open-Volume Corr2	$-\text{corr}(O, V, 10)$
α_7	ADV Conditional	$\begin{cases} -\text{Ts_Rank}(r , 5) & \text{ADV}_{20} < V \\ -1 & \text{else} \end{cases}$
α_8	Intraday Volume	$-\text{rank}(\delta(((C - L) - (H - C))/(H - L) \cdot V, 1))$
α_9	Close Delta	Close delta with min condition
α_{10}	Close Rank	Rank of conditional close delta

C.2 Momentum Alpha Factors

Table 15: Momentum alpha factors.

ID	Name	Formula
α_{11}	1-Week Momentum	$C_t/C_{t-5} - 1$
α_{12}	1-Month Momentum	$C_t/C_{t-20} - 1$
α_{13}	3-Month Momentum	$C_t/C_{t-60} - 1$
α_{14}	6-Month Momentum	$C_t/C_{t-120} - 1$
α_{15}	12-Month Momentum	$C_t/C_{t-252} - 1$
α_{16}	Price Velocity	ΔP_t
α_{17}	Price Acceleration	$\Delta^2 P_t$
α_{18}	Price Jerk	$\Delta^3 P_t$
α_{19}	Volume-Price Trend	$\sum V_i \cdot r_i$
α_{20}	VPT Momentum	$\text{pct_change}(\text{VPT}, 10)$

C.3 Cross-Sectional Alpha Factors

Table 16: Cross-sectional alpha factors.

ID	Name	Formula
α_{21}	A/D Line	$\sum((C - L) - (H - C))/(H - L) \cdot V$
α_{22}	A/D Momentum	$\text{pct_change}(\text{AD}, 10)$
α_{23}	OBV	$\sum \text{sign}(r) \cdot V$
α_{24}	OBV Momentum	$\text{pct_change}(\text{OBV}, 10)$

α_{25}	Breakout High	$\mathbb{1}[C > \max(H, d)], d \in \{10, 20, 50\}$
α_{26}	Breakout Low	$\mathbb{1}[C < \min(L, d)], d \in \{10, 20, 50\}$
α_{27}	Trend Short	$ \text{SMA}_{10} - \text{SMA}_{20} /C$
α_{28}	Trend Long	$ \text{SMA}_{20} - \text{SMA}_{50} /C$

C.4 Statistical Alpha Factors

Table 17: Statistical alpha factors.

ID	Name	Formula
α_{29}	Relative Strength	$C/\text{SMA}_d - 1, d \in \{10, 20, 50\}$
α_{30}	Risk-Adj Momentum	$\bar{r}_d/\sigma_d, d \in \{10, 20\}$
α_{31}	Volume Ratio	$V/\text{SMA}_{20}(V)$
α_{32}	Volatility Rank	$\text{rank}(\sigma_{10})$
α_{33}	Skewness Rank	$\text{rank}(\text{skew}_{20})$
α_{34}	Kurtosis Rank	$\text{rank}(\text{kurt}_{20})$
α_{35}	Hurst Exponent	$H = \log(R/S)/\log(n)$

C.5 High-Frequency Factors

Additional factors for intraday trading:

$$\alpha_{HF1} = \frac{V_t - \text{SMA}_5(V_t)}{\text{SMA}_5(V_t)} \quad (\text{Volume Surge}), \quad (125)$$

$$\alpha_{HF2} = \frac{H_t - L_t}{\text{SMA}_{10}(H - L)} \quad (\text{Range Expansion}), \quad (126)$$

$$\alpha_{HF3} = \frac{C_t - O_t}{H_t - L_t} \quad (\text{Candle Body Ratio}), \quad (127)$$

$$\alpha_{HF4} = \text{corr}(|r_t|, V_t, 20) \quad (\text{Volume-Volatility Correlation}). \quad (128)$$

D Kalman Filter Implementation

The following pseudocode describes the Kalman filter implementation for price denoising:

Algorithm 3 Kalman Filter for Price Denoising

Require: Price series $\{P_t\}_{t=1}^T$, process noise Q , measurement noise R
Ensure: Filtered prices $\{\hat{P}_t\}_{t=1}^T$

- 1: Initialize: $\hat{x}_0 = P_1$, $\hat{P}_0 = 1.0$
- 2: **for** $t = 1$ **to** T **do**
- 3: **Predict:**
- 4: $\hat{x}_{t|t-1} = \hat{x}_{t-1|t-1}$ {Random walk model: $A = 1$ }
- 5: $P_{t|t-1} = P_{t-1|t-1} + Q$
- 6: **Update:**
- 7: $K_t = P_{t|t-1}/(P_{t|t-1} + R)$ {Kalman gain}
- 8: $\hat{x}_{t|t} = \hat{x}_{t|t-1} + K_t(P_t - \hat{x}_{t|t-1})$ {State update}
- 9: $P_{t|t} = (1 - K_t)P_{t|t-1}$ {Covariance update}
- 10: $\hat{P}_t = \hat{x}_{t|t}$
- 11: **end for**
- 12: **return** $\{\hat{P}_t\}$

E Feature Selection Algorithm

Algorithm 4 Robust Feature Calculation with Validation

Require: DataFrame \mathbf{D} with OHLCV columns
Ensure: Feature-enriched DataFrame \mathbf{D}'

- 1: Verify required columns: {close, high, low, volume}
- 2: Convert columns to numeric, coerce errors to NaN
- 3: Drop rows with NaN in essential columns
- 4: **if** $|\mathbf{D}| < 24$ **then**
- 5: **return** \mathbf{D} unchanged
- 6: **end if**
- 7: Compute EMA₁₂, EMA₂₆, MACD, Signal Line
- 8: Compute RSI₁₄ with division-by-zero protection
- 9: Compute ATR₁₄, OBV, Stochastic %K, %D
- 10: Compute lag features (Lag-1 through Lag-3)
- 11: Forward-fill then backward-fill remaining NaN
- 12: **return** \mathbf{D}'

F Database Schema

```

1 CREATE TABLE IF NOT EXISTS positions (
2     id INTEGER PRIMARY KEY AUTOINCREMENT ,
3     session_id TEXT NOT NULL ,
4     symbol TEXT NOT NULL ,
5     side TEXT NOT NULL ,
6     entry_price REAL NOT NULL ,
7     current_price REAL ,

```

```

8   quantity REAL NOT NULL ,
9   investment_amount REAL ,
10  status TEXT DEFAULT 'open' ,
11  ml_confidence REAL ,
12  take_profit_price REAL ,
13  stop_loss_price REAL ,
14  created_at TIMESTAMP DEFAULT CURRENT_TIMESTAMP ,
15  last_update TIMESTAMP DEFAULT CURRENT_TIMESTAMP ,
16  closed_at TIMESTAMP ,
17  exit_price REAL ,
18  pnl REAL
19 );
20
21 CREATE TABLE IF NOT EXISTS trading_sessions (
22  id INTEGER PRIMARY KEY AUTOINCREMENT ,
23  session_id TEXT UNIQUE NOT NULL ,
24  start_time TIMESTAMP DEFAULT CURRENT_TIMESTAMP ,
25  end_time TIMESTAMP ,
26  status TEXT DEFAULT 'active' ,
27  initial_balance REAL ,
28  final_balance REAL
29 );
30
31 CREATE TABLE IF NOT EXISTS ml_predictions (
32  id INTEGER PRIMARY KEY AUTOINCREMENT ,
33  session_id TEXT ,
34  symbol TEXT NOT NULL ,
35  prediction TEXT ,
36  confidence REAL ,
37  predicted_price REAL ,
38  actual_price REAL ,
39  features_used TEXT ,
40  model_version TEXT ,
41  created_at TIMESTAMP DEFAULT CURRENT_TIMESTAMP
42 );

```

Listing 8: Core database schema

G Stacking Ensemble Approach for Trading Signals

Contributed by Devam Patel

G.1 Abstract

This appendix presents a stacking ensemble methodology for generating cryptocurrency trading signals. The approach combines multiple heterogeneous base learners through a meta-learning framework to exploit complementary predictive strengths across different market conditions.

G.2 Base Model Architecture

The stacking ensemble employs five base learners, each capturing different aspects of the price dynamics:

1. **Random Forest:** Captures non-linear feature interactions and provides built-in feature importance ranking. Configuration: 100 trees, max depth 10.
2. **XGBoost:** Sequential boosting with regularization for robust generalization. Configuration: squared error objective, max depth 5.
3. **Linear Regression:** Provides a linear baseline and captures simple trend relationships.
4. **Ridge Regression:** L2-regularized linear model that handles multicollinearity in technical indicators.
5. **Lasso Regression:** L1-regularized model that performs implicit feature selection.

G.3 Meta-Learning Framework

The meta-learner receives the out-of-fold predictions from all base models as input features:

$$\hat{\mathbf{z}}_i = [\hat{y}_i^{(\text{RF})}, \hat{y}_i^{(\text{XGB})}, \hat{y}_i^{(\text{LR})}, \hat{y}_i^{(\text{Ridge})}, \hat{y}_i^{(\text{Lasso})}]^\top \in \mathbb{R}^5. \quad (129)$$

The meta-learner (Gradient Boosting Regressor) learns the optimal combination:

$$\hat{y}_i^{(\text{meta})} = g(\hat{\mathbf{z}}_i; \boldsymbol{\theta}_{\text{meta}}). \quad (130)$$

G.4 Signal Generation

Trading signals are generated by comparing the ensemble prediction to the current price:

$$\text{Signal}_t = \begin{cases} \text{BUY} & \text{if } \hat{P}_{t+1} > P_t \cdot (1 + \tau), \\ \text{SELL} & \text{if } \hat{P}_{t+1} < P_t \cdot (1 - \tau), \\ \text{HOLD} & \text{otherwise,} \end{cases} \quad (131)$$

where τ is a minimum predicted move threshold (typically 0.1–0.5%).

H Mathematical Foundations of Base Models

This appendix provides detailed mathematical derivations for each base model used in the Coinbase trading system.

H.1 Ridge Regression: Derivation

Starting from the penalized least squares objective:

$$J(\boldsymbol{\beta}) = \|\mathbf{y} - \mathbf{X}\boldsymbol{\beta}\|_2^2 + \lambda\|\boldsymbol{\beta}\|_2^2. \quad (132)$$

Taking the gradient and setting to zero:

$$\nabla_{\boldsymbol{\beta}} J = -2\mathbf{X}^\top(\mathbf{y} - \mathbf{X}\boldsymbol{\beta}) + 2\lambda\boldsymbol{\beta} = \mathbf{0}, \quad (133)$$

$$(\mathbf{X}^\top\mathbf{X} + \lambda\mathbf{I})\boldsymbol{\beta} = \mathbf{X}^\top\mathbf{y}, \quad (134)$$

$$\hat{\boldsymbol{\beta}}_{\text{Ridge}} = (\mathbf{X}^\top\mathbf{X} + \lambda\mathbf{I})^{-1}\mathbf{X}^\top\mathbf{y}. \quad (135)$$

The matrix $(\mathbf{X}^\top\mathbf{X} + \lambda\mathbf{I})$ is always invertible for $\lambda > 0$, resolving multicollinearity issues common in technical indicator features.

Bias-Variance Trade-off. As λ increases:

- Bias increases: $\mathbb{E}[\hat{\boldsymbol{\beta}}] \neq \boldsymbol{\beta}$ (biased estimator).
- Variance decreases: $\text{Var}(\hat{\boldsymbol{\beta}}_{\text{Ridge}}) = \sigma^2(\mathbf{X}^\top\mathbf{X} + \lambda\mathbf{I})^{-1}\mathbf{X}^\top\mathbf{X}(\mathbf{X}^\top\mathbf{X} + \lambda\mathbf{I})^{-1}$.

H.2 Lasso Regression: Coordinate Descent

The Lasso objective does not have a closed-form solution due to the non-differentiability of the L1 penalty at zero. The coordinate descent algorithm updates each coefficient:

$$\hat{\beta}_j = \text{sign}(\tilde{\beta}_j) \cdot \max\left(|\tilde{\beta}_j| - \frac{\lambda}{2}, 0\right), \quad (136)$$

where $\tilde{\beta}_j$ is the partial residual regression coefficient:

$$\tilde{\beta}_j = \frac{1}{n} \sum_{i=1}^n x_{ij} \left(y_i - \sum_{k \neq j} x_{ik} \hat{\beta}_k \right). \quad (137)$$

H.3 SVR: Dual Formulation

The dual of the SVR problem is:

$$\max_{\alpha, \alpha^*} -\frac{1}{2} \sum_{i,j} (\alpha_i - \alpha_i^*)(\alpha_j - \alpha_j^*) K(\mathbf{x}_i, \mathbf{x}_j) - \varepsilon \sum_i (\alpha_i + \alpha_i^*) + \sum_i y_i (\alpha_i - \alpha_i^*) \quad (138)$$

$$\text{subject to } \sum_i (\alpha_i - \alpha_i^*) = 0, \quad 0 \leq \alpha_i, \alpha_i^* \leq C.$$

The prediction function is:

$$f(\mathbf{x}) = \sum_{i=1}^n (\alpha_i - \alpha_i^*) K(\mathbf{x}_i, \mathbf{x}) + b. \quad (139)$$

H.4 Gradient Boosting: Functional Gradient Descent

Following ?, gradient boosting performs functional gradient descent in function space.

At iteration m :

1. Compute pseudo-residuals:

$$\tilde{r}_{im} = - \left[\frac{\partial L(y_i, F(\mathbf{x}_i))}{\partial F(\mathbf{x}_i)} \right]_{F=F_{m-1}}. \quad (140)$$

2. Fit a weak learner $h_m(\mathbf{x})$ to pseudo-residuals.

3. Find optimal step size:

$$\rho_m = \arg \min_{\rho} \sum_{i=1}^n L(y_i, F_{m-1}(\mathbf{x}_i) + \rho h_m(\mathbf{x}_i)). \quad (141)$$

4. Update model:

$$F_m(\mathbf{x}) = F_{m-1}(\mathbf{x}) + \eta \cdot \rho_m h_m(\mathbf{x}), \quad (142)$$

where $\eta \in (0, 1]$ is the learning rate (shrinkage parameter).

For squared error loss, the pseudo-residuals are simply the residuals: $\tilde{r}_{im} = y_i - F_{m-1}(\mathbf{x}_i)$.

H.5 XGBoost: Regularized Objective

XGBoost extends gradient boosting with:

1. **Second-order approximation:** Uses both gradient g_i and Hessian h_i for more accurate function approximation.

2. **Structural regularization:** Penalizes model complexity through the number of leaves T and leaf weight magnitude.
3. **Shrinkage:** Scales each tree's contribution by learning rate η .
4. **Column subsampling:** Randomly selects feature subsets at each split, similar to Random Forests.

The optimal structure score for a tree with T leaves is:

$$\tilde{\mathcal{L}} = -\frac{1}{2} \sum_{j=1}^T \frac{\left(\sum_{i \in I_j} g_i \right)^2}{\sum_{i \in I_j} h_i + \lambda} + \gamma T. \quad (143)$$

This score is used in a greedy algorithm to determine the optimal tree structure by evaluating the gain from each candidate split (Equation 65).